OpenCommons: A Distributed, Meaning-Indexed Concept Commons for Accelerating Scientific and Algorithmic Discovery in AI Agents

The Kardashev Project
SI@kardashevproject.org · kardashevproject.org

October 2026

Abstract

When swarms of autonomous AI agents solve complex coding and scientific discovery tasks, they reset their memory between sessions. As a result, agents waste compute repeatedly re-deriving the same mathematical algorithms from scratch while copying subtle numerical errors. Existing skill libraries either store unverified code that spreads bugs across the swarm, or rely on formal verification systems that attempt to translate ordinary Python code into pure mathematical equations—a process that succeeds on only 12.5% of imperative scientific programs. We present **OpenCommons**, a decentralized concept commons that catalogs algorithms by what they actually compute rather than their surface syntax. By executing functions against a small suite of boundary test inputs, OpenCommons identifies equivalent implementations in less than a millisecond, achieving 100% coverage of native scientific code. An asynchronous peer-review engine stress-tests contributed code in the background, quarantining all defective submissions across six adversarial defect classes with zero false alarms. When evaluated on live language models solving scientific discovery tasks, retrieving verified concept cards increases first-attempt success from 33.3% to 83.3% while preventing the propagation of subtle bugs. Across multi-wave research horizons, this verified shared library reduces cumulative discovery cost by 12.67×, providing a dependable foundation for collective machine intelligence.

1 Introduction

When a human scientist or software engineer solves a difficult problem, they package the verified solution into a reusable library so that future projects can build on top of it without starting over. Today's autonomous AI agents, by contrast, largely work in isolation. When asked to build a physics simulation, a data-compression tool, or a statistical analysis pipeline, an AI agent typically writes every helper function from scratch inside its current conversation and forgets what it built the moment the session ends.

Writing foundational code from scratch in every session creates a compounding reliability problem. Because language models generate code by predicting the most common patterns seen in public training data, they naturally default to simplified "textbook" formulas that pass basic sanity checks on toy inputs but silently fail on real-world edge cases. For example, a model asked to compute probabilities will often write a direct exponential sum ($\log \sum_i \exp(x_i)$) that works on small numbers like 1 or 2, but crashes with a numerical overflow as soon as an input reaches 710. Similarly, models routinely write naive summation loops that lose small numbers to computer rounding limits, physics update steps that artificially gain energy over long simulations, or helper functions that accidentally overwrite the caller's original input list in place. When a scientific project requires chaining multiple building blocks together across sequential stages ($W_1, \ldots, W_d$), these subtle failure rates multiply: if an agent has a probability $p_0 < 1$ of writing a bug-free component at any single stage, its probability of completing a d-stage hierarchy without a hidden defect decays exponentially as $\Pr[\text{Pass}(d)] = p_0^d$.

Limitations of Prior Approaches. Prior attempts to give AI agents a shared, long-term memory of reusable code have followed two opposing directions, each hitting a fundamental barrier:

1. **Unverified Code Skill Libraries:** Systems such as Voyager [16], Trove [17], SkillWeaver [15], and shared multi-agent file folders such as CORAL [2] save raw source code or text notes indexed by function names and text descriptions. However, two functions can look completely different on the surface—one written with a `while` loop and another with recursion, or using different variable names—while computing the exact same mathematical result. Because text-based libraries

index code by how it *looks* rather than what it *computes*, agents fill the library with duplicate implementations and conflicting function names. Worse, unverified libraries suffer from *epistemic contagion*: if Agent A saves a function that passes a toy test but contains a hidden edge-case bug (such as dividing by zero on empty inputs or mutating an input list in place), Agent B will later retrieve and trust that broken function, spreading the failure across the entire group of agents. At scale, as the library grows to $N = 5{,}000$ functions, scanning flat files consumes 61,800 tokens per search, while standard keyword search (BM25) confuses closely named algorithms (`base58` vs. `base64`, `crc16` vs. `crc32`, `rk4` vs. `rk2`), dropping top-1 search accuracy to 32.5%.

2. **Formal IR Lifting and Theorem Proving:** At the opposite extreme, formal verification systems attempt to mathematically prove that contributed code is correct before storing it. Because automated theorem provers (such as Z3 [3] and Lean 4 [9]) cannot easily reason about ordinary imperative Python code that modifies variables and arrays in place, these systems first try to translate ("lift") the Python code into a restricted set of pure mathematical equations without mutable variables—a representation known as a *functional intermediate representation* (IR). In practice, this hits the **Formal Lifting Expressiveness Wall**: 87.5% (210/240) of real-world scientific and systems Python functions cannot be translated into pure equations because they rely on standard loops, running state accumulators, byte buffers (`bytearray`), and hash tables. Across our 3-Repository Multi-File Engineering Suite, functional lifting rejected 83.3%–91.8% of functions, while synchronous theorem proving succeeded on only 39.3% of cross-algorithm equivalence checks, stalled for 1.53 s per failed proof search, and added a +6,433-token/turn tool-schema overhead that inflated empty-library token costs to 1.399×.

System Architecture: OpenCommons. We introduce **OpenCommons**, a decentralized, behavior-indexed Concept Commons that bridges these two extremes: it accepts 100% of native Python code without requiring translation into pure math equations, while still guaranteeing that stored code is behaviorally deduplicated and stress-tested against edge-case bugs. OpenCommons is built on four connected mechanisms:

- **1. Indexing Code by Behavior (sid + fp → cid):** Instead of asking a slow theorem prover whether two functions are equivalent, OpenCommons simply runs each function on a standardized set of 16 test inputs ("probes")—covering zeroes, negative numbers, primes, and extreme boundary values—and hashes the resulting outputs into a behavioral fingerprint (fp, $\approx$ 0.33 ms). Combined with a fast structural syntax-tree check (sid, $\approx$ 0.15 ms) that strips away variable-name differences and extracts branch threshold constants, OpenCommons groups all equivalent algorithms under a single canonical concept ID (cid) in 0.48 ms.

- **2. Fast Retrieval with Live Input-Output Verification:** Agents search the library through a lightweight two-tool interface (`find` and `record`). To prevent keyword confusion between similar algorithms (like `crc16` and `crc32`), `find` checks both the exact algorithm number and the function's type signature, and optionally runs candidate functions live in under 2 ms against a sample input-output pair provided by the agent. It then returns the verified function together with any helper functions it depends on.

- **3. Background Stress-Testing & Lock-Free Sync (CRDTs):** To catch hidden bugs without slowing down the agent, an asynchronous peer-review worker tests every newly contributed function in the background against five automated verification gates (checking for constant placeholder stubs, wrong protocol constants, accidental in-place list mutations, NaN/infinity overflows, and zero-division edge cases). Any function that fails is automatically quarantined so other agents never use it. When multiple agents work in parallel across different machines, they merge their libraries using a Conflict-Free Replicated Data Type (CRDT)—a mathematical merge rule that guarantees every machine converges to the exact same verified library without a central coordinator.

- **4. Dual Concept Cards to Prevent Blind Copying:** When an agent retrieves a stored function to solve a modified problem, showing only raw Python code often causes *syntactic copy-bias*: the language model blindly copies the old variable names and lines verbatim instead of adapting the logic to the new task. OpenCommons therefore serves each concept as a *Dual Concept Card* that pairs a concise step-by-step pseudocode summary, physical conservation laws and units, and explicit "what-not-to-do" pitfall warnings alongside the verified reference code.

Summary of Contributions. Every claim in this paper is backed by pre-registered, confound-audited experiments executed under our verification-first evaluation harness:

1. **Formal Model & CRDT/Horizon Theorems**

(Section 3): We formalize sid and fp/cid (Definitions 1–2), prove Strong Eventual Consistency of the federated store $(\mathcal{L}, \sqcup)$ (Theorem 1), derive the Compositional Horizon Scaling Law showing exponential cost growth $\Theta(d \cdot \bar{L}_{\text{syn}} \cdot p_0^{-d})$ for isolated agents vs. $O(\bar{L}_{\text{syn}} + d \cdot \bar{L}_{\text{card}})$ for OpenCommons (Theorem 2), and establish an information-theoretic conditioned generation model explaining why pseudocode and pitfall cards outperform raw code on adaptation tasks (Proposition 1).

2. **Breaking the Formal Lifting Wall (Section 5.1):** Across 240 scientific and systems functions, 120 cross-algorithm equivalence pairs, and 120 subtle semantic mutants, OpenCommons achieves 100.0% native ingestion coverage (8.0× over formal lifting's 12.5%), 0.48 ms mean indexing latency, 100.0% cross-algorithm equivalence recall, and 0.0% false merges.

3. **Sub-10ms Retrieval Scaling to $N = 5{,}000$ & Modality Compression (Section 5.2):** Across $N \in \{120, 500, 1{,}000, 2{,}500, 5{,}000\}$ concepts, OpenCommons sustains 100.0% Precision@1 at 7.24 ms mean latency (+67.5 pts over BM25 and 650.5× fewer scanned tokens than CORAL), while pseudocode building blocks achieve 2.67× token compression over raw code.

4. **100% Defect Quarantine Across 6 Classes & CRDT Convergence (Section 5.3):** Asynchronous peer review quarantines 100.0% (150/150) of subtle defective implementations across 6 defect classes with 0.0% (0/50) false quarantines, and 10-node CRDT federation achieves 100.0% order-invariance across 50 random merge permutations.

5. **Live LLM Scientific Discovery & Multi-Wave Horizon Scaling (Sections 5.4–5.5):** On live `Qwen3.5-2B-Base` evaluation across 30 scientific tasks, Dual Concept Cards boost Pass@1 by +50.0 percentage points over isolated zero-shot generation (83.3% vs. 33.3%), outperform raw code on compositional adaptation by +25.0 pts (75.0% vs. 50.0%), and shift exact log-probability preference margins toward canonical implementations by +0.165 nats/token. Across 5 compositional waves and 234 agent sessions, OpenCommons eliminates empty-library carrying taxes (E/A = 0.907×), reduces cross-project convention cost-to-pass to 0.419×, and achieves a 12.67× Wave-5 cost reduction over isolated agents (6.40× over CORAL).

6. **Adversarial Stress-Testing & Empirical Boundary Analysis (Section 5.6):** We probe empirical failure boundaries: probe-count sensitivity ($K = 16 \rightarrow 0.0\%$ collisions), floating-point ULP drift stability up to condition number $\kappa = 10^7$ (subtractive cancellation breaks equivalence at $\kappa \geq 10^8$), 20-node CRDT convergence under 50% packet loss (100/100 schedules), live LLM epistemic contagion from subtle bugs (13.3% vs. 83.3%), white-box AST comparator boundary seeding (100% → 0% false merges on threshold branches), and token-level proof of syntactic copy-bias (+0.794 nats/token net disentanglement).

2 Related Work and Novelty Positioning

The challenge of enabling autonomous AI agents to accumulate, verify, and reuse algorithmic knowledge across tasks sits at the intersection of two distinct research traditions: *LLM agent skill libraries* and *formal program verification and library learning*. Table 1 contrasts OpenCommons with representative systems across both areas.

2.1 Agent Skill Libraries and Shared Memory Systems

Recent agentic frameworks equip language models with external memory stores to persist reusable skills across episodes. Voyager [16] introduced a skill library of JavaScript programs for embodied exploration, indexed by natural-language description embeddings. Trove [17] induces reusable Python toolboxes from problem-solving trajectories via periodic syntactic pruning, while SkillWeaver [15] synthesizes compositional API wrappers for web agents. CORAL [2] introduced a shared multi-agent workspace where autonomous agents read and write flat Markdown skill files and notes across collaborative tasks.

While these systems demonstrate that persisting code across sessions can improve downstream task completion, they exhibit four core limitations:

1. **Absence of Behavioral Equivalence Classes:** Two functions can look completely different on the surface—such as an iterative while loop versus tail recursion, or different variable names—yet compute the exact same mathematical transformation from inputs to outputs. Existing skill libraries index artifacts solely by surface identifier names, docstrings, or dense text embeddings. Because they lack a formal notion of *behavioral equivalence*, independent agents that discover the same underlying algorithm under different variable names or loop structures proliferate redundant entries. Distinct agents registering incompatible behavioral parameterizations under the same function name trigger naming collisions.

Table 1: **Architectural comparison of OpenCommons against prior agent skill libraries and formal verification frameworks.** Unlike unverified text/code repositories that suffer from keyword collisions and silent bug propagation, or formal IR lifters that reject $> 80\%$ of imperative scientific Python, OpenCommons achieves 100% native-code coverage with sub-millisecond behavioral equivalence indexing, live I/O-probe retrieval, asynchronous adversarial quarantine, and lock-free CRDT federation.

System / Architecture	Native Imperative Coverage	Behavioral Equiv. Classes (cid)	Live I/O Probe Retrieval (< 8 ms)	Async Defect Quarantine	Lock-Free CRDT Federation	Modality Ablation	Empirical Latency
Voyager [16]	✓	× (Embedding)	×	×	× (Single-agent)	× (Raw code)	50–200 ms
Trove [17]	✓	× (Syntactic)	×	×	× (Single-agent)	× (Raw code)	10–100 ms
SkillWeaver [15]	✓	× (API schema)	×	×	× (Centralized)	× (Raw code)	20–150 ms
CORAL [2]	✓	× (Flat Markdown)	× ($O(N)$ grep)	×	× (Filesystem)	× (Text notes)	$O(N)$ turns
DreamCoder / LILO [4, 5]	× (Pure λ-IR)	× (Syntactic β-eq)	×	×	×	× (DSL primitives)	10^3–10^5 ms
Verus / Dafny / Lean 4 [8, 7, 9]	× (Proof DSL)	✓(SMT / Type)	×	✓(Sync prover)	×	× (Formal spec)	0.24–15 s
Formal-IR-Lift (21-Form Pure Functional)	12.5%	Partial (39.3%)	×	Sync Z3/Lean	×	× (Raw / IR)	244–1,530 ms
OpenCommons (Ours)	**100.0%**	**✓ (sid + 16-probe fp)**	**✓ (100% P@1)**	**✓ (100% recall)**	**✓ (G-Set ⊔)**	**✓ (4 Modalities)**	**0.48 ms ingest**

2. **Linear Search Overhead and Near-Miss Retrieval Pollution:** In flat-file workspaces such as CORAL, agents execute $O(N)$ filesystem scans (`grep`/`find`) to locate prior code, consuming 61,800 scanned tokens at $N = 5{,}000$ concepts and adding +3.8–4.4 turns per wave. Standard lexical retrieval (BM25) over code repositories conflates numbered algorithmic families (`base58` vs. `base64`, `crc16` vs. `crc32`, `rk4` vs. `euler`), degrading Precision@1 to 32.5% at $N = 5{,}000$. These systems omit sub-millisecond live input-output probes during retrieval to guarantee that returned candidates satisfy caller semantics.

3. **Unverified Defect Propagation (Epistemic Contagion):** If Agent A implements a function that passes a simple test case but has a hidden bug on edge cases (such as zero division, unhandled boundary conditions, or silently mutating caller input containers in place), Agent A will save that code to the unverified library. Downstream Agent B retrieves and trusts that code, causing *epistemic contagion* where subtle defects propagate across the swarm and cause cascading multi-wave failures.

4. **Representation Modalities and Syntactic Copy-Bias:** Prior toolboxes uniformly serve retrieved skills as raw source code or unstructured text notes. When a downstream task requires adapting an algorithm to modified inputs or constraints, providing raw code induces *syntactic copy-bias*: the language model's attention mechanism latches onto exact variable names and operators in the reference code and copies them verbatim rather than adapting the underlying logic. Prior systems do not investigate how structured concept representations (compact algorithmic pseudocode, mathematical invariants, and refuted pitfall warnings) can prevent this copy-bias and improve adaptation accuracy.

2.2 Formal Verification, Program Synthesis, and the Expressiveness Wall

A separate line of research seeks mathematical guarantees through formal program verification and neuro-symbolic library learning. To prove that a program is mathematically correct, traditional verification frameworks require programs to be expressed in pure functional representations without mutable variables, such as λ-calculi or proof-annotated domain-specific languages (DSLs) [4, 1, 5, 8, 7, 9]. In compiler systems, translating ordinary imperative source code into such a pure mathematical formalism is known as *intermediate representation (IR) lifting*. Ordinary scientific Python is written imperatively: code updates variables in place, accumulates values across loops, and mutates memory buffers. An IR lifter attempts to convert this mutable control flow into a clean set of mathematical equations where no variable is ever reassigned, allowing automated theorem provers and SMT solvers to verify equivalence.

Grammar of Formal-IR-Lift. To evaluate the limits of formal intermediate representations, we formulated Formal-IR-Lift, a baseline that lifts imperative Python into an explicit 21-form pure functional grammar. Let τ denote first-order monomorphic types and e denote expressions:

$$
\begin{aligned}
\tau ::= {} & \mathrm{Int} \mid \mathrm{Bool} \mid \mathrm{String} \mid \mathrm{Float} \mid \mathrm{Rat} \mid \mathrm{Word}_w \mid \mathrm{List}[\tau] \\
& \mid \mathrm{Option}[\tau] \mid \tau_1 \times \tau_2 \mid \tau_1 \to \tau_2 \mid \mathrm{Record}\{\ell_i : \tau_i\}, \\
e ::= {} & x \mid c \mid \lambda x : \tau.\, e \mid e_1\, e_2 \mid \mathrm{let}\ x = e_1\ \mathrm{in}\ e_2 \\
& \mid \mathrm{if}\ e_1\ \mathrm{then}\ e_2\ \mathrm{else}\ e_3 \mid e_1 \oplus e_2 \mid e_1 \odot e_2 \mid \neg e \\
& \mid e_1 \wedge e_2 \mid e_1 \vee e_2 \mid \mathrm{nil} \mid \mathrm{cons}(e_1, e_2) \\
& \mid \mathrm{foldr}(f, e_0, e_{\mathrm{list}}) \mid \mathrm{natrec}(n, e_0, f) \mid \mathrm{loop}(n, s_0, f) \\
& \mid \mathrm{pair}(e_1, e_2) \mid \mathrm{fst}(e) \mid \mathrm{snd}(e) \mid \mathrm{none} \\
& \mid \mathrm{some}(e) \mid \mathrm{case_opt}(e, e_{\mathrm{none}}, f_{\mathrm{some}}) \mid \{\ell_i = e_i\} \mid e.\ell.
\end{aligned}
\tag{1}
$$

The Expressiveness Wall. Evaluating Formal-IR-Lift against real-world engineering code exposed three fundamental bottlenecks:

1. **Domain Incompatibility with Imperative Scientific Code:** Real-world scientific computing, numerical simulation, cryptography, and systems engineering rely on imperative loops (`while`/`for`), multi-variable state updates (Kahan compensation, Welford online moments, Runge–Kutta stages), in-place byte buffers (`bytearray`), hash tables (`dict`/`set`), exception handling, and domain libraries (`math`, `numpy`, `re`). Pure functional IR lifting accepts only 12.5% (30/240, 95% CI [8.8%, 16.7%]) of standard scientific and systems functions—meaning 87.5% of real-world scientific code cannot be translated into pure functional equations without mutable state. We validated this wall across our **3-Repository Multi-File Engineering Suite**:
   - *Neural Language Model Training Suite*: A 90-function deep learning training framework featuring multi-head attention, RoPE positional embeddings, fused layer normalization, and dynamic learning-rate schedules. Evaluated on learning-rate schedule synthesis, functional lifting rejected 83.3% of functions (75/90).
   - *Distributed Message-Routing Service*: A 147-function asynchronous event broker implementing priority queuing, protocol framing, dead-letter routing, and dynamic ranking. Evaluated on ranking pipeline refactoring, functional lifting rejected 91.8% of functions (135/147).
   - *Systems Benchmark Suite*: A 149-function systems benchmark featuring sliding-window statistics, circular ring buffers, and streaming deduplication. Evaluated on streaming deduplication, functional lifting rejected 89.9% of functions (134/149).
2. **Inductive Bottlenecks on Cross-Algorithm Equivalence:** Even on programs that lift into functional IR, proving equivalence between distinct algorithmic structures (Euclidean vs. binary Stein GCD, or iterative accumulator loops vs. tail-recursive definitions) requires synthesizing inductive loop invariants. Z3 and Lean 4 plateaued at 39.3% (118/300) recall on equivalence pairs, while syntactic anti-unification in Stitch achieved −2.52% description length compression and 1.2% downstream reuse.
3. **Synchronous Prover Latency and Context-Window Tax:** Placing SMT solvers and theorem provers on the synchronous agent critical path caused CPU stalls, with Lean 4 spending 90% of compute on failed proof searches (1.53 s mean failure latency vs. 244 ms on trivial successes). Exposing an 18-tool formal verification MCP schema imposed a +6,433-token per-turn context tax, inflating empty-library session token costs to 1.399× the no-tool baseline.

OpenCommons resolves this dichotomy by decoupling fast online behavioral indexing from asynchronous background refutation. By operating directly on native Python ASTs and executing stratified canonical input probes in < 0.5 ms, OpenCommons achieves 100.0% ingestion coverage (8.0× higher than formal IR lifting, $p = 0.0001$) and 100.0% cross-algorithm equivalence recall, while shifting adversarial boundary fuzzing and contract verification off the agent's critical path into a lock-free, join-semilattice CRDT store.

3 OpenCommons Architecture & Formal Mathematical Model

We formalize the core data structures, equivalence relations, distributed convergence guarantees, and information-theoretic scaling laws of OpenCommons. Let $\mathcal{P}$ denote the space of syntactically valid native Python function definitions $f : \mathcal{X}_1 \times \cdots \times \mathcal{X}_k \to \mathcal{Y}$ annotated with parameter type signatures $\tau_{\text{in}}(f) = (T_1, \ldots, T_k)$ and return type $\tau_{\text{out}}(f) = T_{\text{ret}}$.

3.1 Two-Tier Meaning-Addressed Identity (sid, fp, and cid)

Two functions can look completely different on the surface—for example, computing the greatest common divisor using an iterative while loop versus recursion, or using different variable names—yet compute the exact same mathematical mapping from inputs to outputs. To catalog code by meaning rather than surface appearance, a system must identify this *behavioral equivalence*. Proving behavioral equivalence symbolically for general programs is undecidable by Rice's Theorem [11], and theorem provers stall when unrolling complex loops. Instead of slow symbolic theorem proving, OpenCommons evaluates each function on a carefully chosen set of test inputs ("probes")—including boundary edge cases like zeroes, negative numbers, and extreme dynamic ranges—and hashes its observed outputs to produce a compact behavioral fingerprint. Two functions with identical fingerprints belong to the same canonical equivalence class (cid). To achieve both syntax-invariant structural deduplication and behavioral equivalence without restricting language expressiveness, OpenCommons assigns every ingested function $f \in \mathcal{P}$ a two-tier identity: an alpha-normalized structural Merkle hash (sid) and a canonical execution trajectory fingerprint (fp).

Definition 1 (Structural Merkle Identity (sid))**.** *Let $AST(f) = (V, E, \ell)$ be the parsed abstract syntax tree of*

function $f \in \mathcal{P}$ with $|V| = |T|$ nodes. Define the alpha-commutative normalization operator $\mathcal{N} : \mathcal{T} \to \mathcal{T}$ via a deterministic top-down/bottom-up tree traversal:

1. ***Docstring and Decorator Stripping:*** *Remove non-executable string literal headers, type annotations, and decorator nodes.*
2. ***Positional De Bruijn Alpha-Renaming:*** *Map the function's self-identifier to $\$self$, the i-th formal parameter to $\$p_i$ ($i \in \{0, \dots, k-1\}$), and the j-th local variable bound in store context in pre-order traversal to $\$v_j$ ($j \in \mathbb{N}_0$), substituting load occurrences accordingly. Normalize annotated assignments to untyped assignments.*
3. ***Commutative Subtree Canonicalization:*** *For every binary operation node $u = BinOp(\oplus, L, R)$ where $\oplus \in \{+, \times, \&, |, \wedge\}$ or comparison node $u = Compare(\odot, L, R)$ where $\odot \in \{==, \neq\}$, recursively normalize $L' = \mathcal{N}(L)$ and $R' = \mathcal{N}(R)$, and order the children lexicographically by their serialized representations $dump(L') \leq_{lex} dump(R')$.*

*The **Structural Merkle Identity** $h_{sid} : \mathcal{P} \to \{0,1\}^{64}$ is the 64-bit truncated SHA-256 digest of the canonical AST serialization:*

$$h_{sid}(f) \triangleq \mathrm{Trunc}_{64}\Big(\mathrm{SHA256}\left(\mathrm{serialize}(\mathcal{N}(AST(f)))\right)\Big). \tag{2}$$

Because $\mathcal{N}$ visits each AST node a constant number of times and performs bounded lexicographical comparisons on serialized subtrees, $h_{sid}(f)$ runs in $O(|T|)$ time (≈ 0.15 ms empirically) and satisfies $h_{sid}(f_1) = h_{sid}(f_2)$ for all α-equivalent and commutatively reordered implementations. Tree traversal extracts all called function identifiers $deps(f) = (g_1, \dots, g_m)$ to support transitive dependency bundle closure.

While h_{sid} collapses variable renamings, comment variations, and commutative operand swaps, distinct algorithms computing the same mathematical function (iterative Euclidean GCD vs. tail-recursive GCD) possess non-isomorphic ASTs. To unify cross-algorithm implementations into a single semantic concept, we introduce the canonical behavioral fingerprint.

Definition 2 (Canonical Behavioral Fingerprint (fp) and Equivalence Class (cid))**.** *For parameter signature $\tau_{in} = (T_1, \dots, T_k)$, let $\mathcal{G}(\tau_{in}, K) = \{\mathbf{x}_1, \dots, \mathbf{x}_K\}$ be a deterministic, stratified suite of $K = 16$ canonical probe tuples covering boundary values, zero/identity elements, negative/fractional regimes, and adversarial boundary inputs for each type T_j (`int`, `float`, `bool`, `str`, `bytes`, `bytearray`, `list[float]`, `list[int]`, `dict`, `tuple`), with parameter index offset $\phi(i,j) = (i + 3j) \bmod 16$ for $i \geq 4$ to prevent diagonal aliasing ($x_{i,1} = x_{i,2}$). To cover extreme numerical regimes, $\mathcal{G}$ includes 8 adversarial boundary probes: dynamic range overflow ($[1000.0, 1001.0, 1002.0]$), underflow ($[-1000.0, -1001.0, -1002.0]$), large magnitude ($[10^{300}, 10^{300}]$), degenerate zero-variance ($[0.0, 0.0, 0.0]$), singleton variance singularity ($[42.0]$), empty boundary ($[]$), subnormal floats ($[10^{-300}, -10^{-300}]$), and subtractive cancellation ($[10^{15}, 1.0, -10^{15}]$).*

Let $\sigma : \mathcal{V} \to \Sigma^*$ be a deterministic value canonicalizer that maps integers, booleans, strings, and byte sequences to exact representations, sorts unordered containers (`set`, `dict`), and quantizes IEEE-754 floating-point numbers to 12 significant decimal digits with zero-clamping below $\epsilon_0 = 10^{-12}$:

$$\sigma_{float}(z) = \begin{cases} 0 & \text{if } |z| < 10^{-12}, \\ \mathrm{round}_{12g}(z) & \text{if } |z| \geq 10^{-12}, \\ NaN \mid \pm\infty & \text{otherwise.} \end{cases} \tag{3}$$

For each probe $\mathbf{x}_i \in \mathcal{G}(\tau_{in}, K)$, we execute f on a deep copy $\tilde{\mathbf{x}}_i = \mathrm{clone}(\mathbf{x}_i)$ and define the argument-mutation detector $\Delta_{mut}(\mathbf{x}_i) \triangleq \mathbb{I}[\sigma(\tilde{\mathbf{x}}_i^{post}) \neq \sigma(\mathbf{x}_i^{pre})]$ and the probe outcome token $o_i(f)$:

$$o_i(f) \triangleq \begin{cases} \texttt{OK:}\sigma(f(\tilde{\mathbf{x}}_i))\texttt{:M(}\sigma(\tilde{\mathbf{x}}_i^{post})\texttt{)} & \text{if } \Delta_{mut}(\mathbf{x}_i) = 1, \\ \texttt{OK:}\sigma(f(\tilde{\mathbf{x}}_i))\texttt{:M()} & \text{if } \Delta_{mut}(\mathbf{x}_i) = 0, \\ \texttt{ERR:}\,\mathrm{type}(e) & \text{on exception } e. \end{cases} \tag{4}$$

*Define the constant-stub indicator $\chi_{const}(f) \triangleq \mathbb{I}\big[|\{i : o_i(f) \text{ is } \texttt{OK}\}| \geq 4 \ \wedge \ |\{\sigma(f(\tilde{\mathbf{x}}_i))\}| = 1 \ \wedge \ \tau_{out} \notin \{\texttt{bool}, \texttt{None}\} \ \wedge \ k \geq 1\big]$. The **Canonical Behavioral Fingerprint** $h_{fp} : \mathcal{P} \to \{0,1\}^{64}$ is:*

$$h_{fp}(f) \triangleq \mathrm{Trunc}_{64}\Big(\mathrm{SHA256}\Big(\big\|_{i=1}^{K} o_i(f)\Big)\Big), \tag{5}$$

*with fallback $h_{fp}(f) = h_{sid}(f)$ only if all K generic probes raise domain exceptions. For piecewise threshold functions (such as Reynolds transition $Re > 2300$), white-box AST inspection extracts comparator numeric literals c and injects boundary probes $\{c - \epsilon, c, c + \epsilon\}$ to eliminate threshold blindness. Two functions $f_1, f_2 \in \mathcal{P}$ are behaviorally equivalent, written $f_1 \sim f_2$, if and only if $h_{fp}(f_1) = h_{fp}(f_2)$. The quotient space $\mathcal{P}/\sim$ partitions programs into **Canonical Equivalence Classes** $[f]_\sim$ identified by $cid(f) \triangleq \texttt{cid:}h_{fp}(f)$ in ≈ 0.33 ms (0.48 ms combined sid + fp latency).*

Remark 1 (Avoidance of Rice's Theorem and Floating-Point Reassociation Drift)**.** *Exact program equivalence across infinite domains is undecidable [11], and bounded SMT translation times out when unrolling loops with non-linear arithmetic. Definition 2 replaces symbolic decision procedures with observational equivalence over a stratified K-point finite algebra, where σ_{float} absorbs $O(N\epsilon_{mach})$ IEEE-754 non-associativity between algebraically equivalent evaluation orders while separating unstable approximations on boundary probes.*

3.2 Asynchronous 5-Gate Peer Review and Join-Semilattice CRDT Federation

In a collaborative swarm of discovery agents, a key vulnerability is *epistemic contagion*: if Agent A implements a function that passes a simple test case but has a hidden bug on edge cases (such as zero division, unhandled boundary conditions, or mutating an input list in place), Agent A will save that code to the shared library. Later, Agent B retrieves and trusts that code, spreading the defect across the swarm and causing cascading failures. Running exhaustive verification synchronously on the agent's critical path would stall discovery. OpenCommons resolves this dilemma through *asynchronous peer review*: newly contributed functions are immediately accessible locally, while background workers stress-test them on adversarial edge cases and quarantine any broken implementations before other agents can be infected.

When multiple agents work in parallel on different machines without a central lock or database coordinator, they need a way to merge their shared libraries so everyone ends up with the exact same library regardless of network packet drops, delays, or out-of-order message arrival. We achieve this by formalizing the store as a Conflict-Free Replicated Data Type (CRDT) [14]. In a CRDT, the merge operation between two replicas is mathematically defined to be commutative, associative, and idempotent, ensuring that all agent nodes eventually converge to the identical verified library state without centralized coordination.

The 5-Gate Asynchronous Adversarial Peer-Review Engine. When a block b is ingested, it is immediately assigned trust tier $\tau_{\text{tier}}(b) = \texttt{probe_verified}$ in < 0.5 ms. Off the agent's critical path, the background engine audits b across five refutation gates:

- **Gate 1 (Zero-Valid-Probe Guard):** Quarantines b if all $K = 16$ canonical probes raise runtime exceptions and no contract I/O assertions are supplied.
- **Gate 2 (Degenerate Constant-Stub Refutation):** Quarantines b if $\chi_{\text{const}}(b) = 1$ (returning a constant literal across ≥ 4 distinct valid probe inputs on a non-boolean return type).
- **Gate 3 (Silent Caller-Argument Mutation Guard):** Quarantines b if $\Delta_{\text{mut}}(\mathbf{x}_i) = 1$ on any probe for a value-returning function unless in-place mutation is declared.
- **Gate 4 (Contract I/O Verification):** Executes b on deep-copied caller/domain contract pairs $\{(\mathbf{x}_c, y_c)\}$, quarantining b on output discrepancy or exception (catching wrong protocol seeds such as 0x0000 vs. 0xFFFF).
- **Gate 5 (Extended 32-Probe Differential & Invariant Fuzzing):** Evaluates b over an extended 32-probe boundary manifold equipped with the 8 adversarial boundary probes against mathematical domain invariants $\mathcal{I}(\mathbf{x}, f(\mathbf{x}))$ and differential reference implementations.

Blocks passing all five gates are promoted to $\tau_{\text{tier}}(b) = \texttt{peer_reviewed}$; blocks failing any gate are demoted to `quarantined` with an explicit counterexample trace.

State-Based Join-Semilattice CRDT Store. We formalize the distributed store as a state-based Conflict-Free Replicated Data Type (CvRDT) [14]. Let $\mathcal{B}$ be the universe of immutable `ConceptBlock` records $b = (\text{id}, \text{cid}, \text{sid}, \text{fp}, \tau_{\text{tier}}, \text{pitfalls}, \dots)$, where $\text{id}(b) = \texttt{name}(b)\texttt{@}\text{sid}(b)_{1:8}$ is content-addressed. Each block carries a verification trust tier $\tau_{\text{tier}}(b) \in \{\texttt{quarantined} \prec \texttt{probe_verified} \prec \texttt{peer_reviewed}\}$. Define the total quality ordering $\prec_q$ on blocks sharing concept cid by lexicographically comparing:

$$\tau(b) \triangleq \Big(\mathbb{I}[\tau_{\text{tier}}(b) = \texttt{peer_reviewed}], \; |\text{pitfalls}(b)|, \; -\text{id}(b)_{\text{lex}}\Big), \tag{6}$$

where $-\text{id}(b)_{\text{lex}}$ provides a deterministic tie-breaker via smallest content-addressed identifier.

Theorem 1 (Strong Eventual Consistency of Federated Concept Commons)**.** *Define the state of node m in OpenCommons as a triple $\mathcal{S} = (B, Q, C)$, where:*

- *$B \subseteq \mathcal{B}$ is a grow-only set (G-Set) of contributed concept blocks indexed by id(b),*
- *$Q \subseteq \{id(b) : b \in B\}$ is a monotone quarantine set recording blocks refuted by asynchronous peer review (Gates 1–5), and*
- *$C : CID \rightharpoonup B$ is the canonical representative selector defined for each active equivalence class cid with $B_{cid} \setminus Q \neq \emptyset$ by:*

$$C(cid) \triangleq \operatorname*{arg\,max}_{b \in B_{cid} \setminus Q} \tau(b). \tag{7}$$

Define the binary state-merge operator $\sqcup : \mathcal{L} \times \mathcal{L} \to \mathcal{L}$ between two store replicas $\mathcal{S}_1 = (B_1, Q_1, C_1)$ and $\mathcal{S}_2 = (B_2, Q_2, C_2)$ as:

$$\mathcal{S}_1 \sqcup \mathcal{S}_2 \triangleq \Big(B_1 \cup B_2, \; Q_1 \cup Q_2, \; C_{1\sqcup 2}\Big), \tag{8}$$

where $C_{1\sqcup 2}(cid) = \arg\max_{b \in (B_1 \cup B_2)_{cid} \setminus (Q_1 \cup Q_2)} \tau(b)$ (and for any id $\in B_1 \cap B_2$, the block record promotes τ_{tier} to `quarantined` if id $\in Q_1 \cup Q_2$, or

to $\max(\tau_{tier}(b_1), \tau_{tier}(b_2))$ *otherwise). Then* $(\mathcal{L}, \sqcup, \mathcal{S}_\perp)$ *with* $\mathcal{S}_\perp = (\emptyset, \emptyset, \emptyset)$ *forms a bounded join-semilattice satisfying Associativity, Commutativity, and Idempotence (ACI), guaranteeing lock-free Strong Eventual Consistency (SEC) across arbitrary network topologies and message delivery permutations.*

Proof. Because $C(\text{cid})$ is a deterministic pure projection of (B, Q), the state space $\mathcal{L}$ is isomorphic to the product lattice $\mathcal{P}_{\text{fin}}(\mathcal{B}) \times \mathcal{P}_{\text{fin}}(\text{ID})$ ordered by componentwise set inclusion $(B_1, Q_1) \sqsubseteq (B_2, Q_2) \iff B_1 \subseteq B_2 \ \wedge \ Q_1 \subseteq Q_2$, with bottom element $\mathcal{S}_\perp = (\emptyset, \emptyset, \emptyset)$.

1. **ACI of** (B, Q)**:** Set union $\cup$ on $\mathcal{P}_{\text{fin}}(\mathcal{B})$ and $\mathcal{P}_{\text{fin}}(\text{ID})$ is associative ($(A_1 \cup A_2) \cup A_3 = A_1 \cup (A_2 \cup A_3)$), commutative ($A_1 \cup A_2 = A_2 \cup A_1$), and idempotent ($A_1 \cup A_1 = A_1$). When two replicas hold the same id with different local trust tiers (e.g., replica 1 holds `probe_verified` while replica 2 has placed $\text{id} \in Q_2$), the quarantine rule $\text{id} \in Q_1 \cup Q_2 \implies \tau_{\text{tier}} =$ `quarantined` acts as a monotone absorption join over `probe_verified` $\leq$ `peer_reviewed` $\leq$ `quarantined`.

2. **Uniqueness and Order-Invariance of** $C_{1 \sqcup 2}(\textbf{cid})$**:** For fixed (B, Q) where $B = \bigcup_{m=1}^{M} B_m$ and $Q = \bigcup_{m=1}^{M} Q_m$, the candidate set $A_{\text{cid}} = \{b \in B : \text{cid}(b) = \text{cid} \ \wedge \ \text{id}(b) \notin Q\}$ is invariant to replica merge order. In the CRDT store merge operator, blocks are canonicalized by sorting over $(\text{block_id}, \text{cid})$ prior to insertion, and $\tau(b)$ defines a strict total order on A_{cid} whenever $\text{id}(b)$ is distinct. Therefore, the maximum $\arg\max_{b \in A_{\text{cid}}} \tau(b)$ is unique and independent of merge order.

3. **Monotonicity of Local Operations:** Local ingestion adds a block b to B ($B \to B \cup \{b\}$), and asynchronous review either promotes $\tau_{\text{tier}}(b)$ to `peer_reviewed` or inserts $\text{id}(b)$ into Q ($Q \to Q \cup \{\text{id}(b)\}$). Every local transition is inflationary with respect to $\sqsubseteq$.

By the fundamental theorem of state-based CRDTs [14], any set of M nodes that exchange state updates over a connected gossip graph converges to the exact least upper bound $\bigsqcup_{m=1}^{M} \mathcal{S}_m$ without coordination locks. □

3.3 Compositional Horizon Scaling and Minimum Description Length Advantage

We analyze how expected token expenditure scales as scientific discovery tasks compose deeper hierarchies of previously discovered concepts across multi-wave horizons $W_1, \ldots, W_d$, and how a minimal 2-tool MCP interface (`find`, `record`) eliminates the carrying cost of legacy 18-tool formal schemas.

Theorem 2 (Compositional Horizon Scaling Law)**.** *Consider a d-wave compositional scientific discovery benchmark where solving a target task at wave depth* $d \in \mathbb{N}^+$ *requires synthesizing a top-level function* $f^{(d)}$ *that depends transitively on a hierarchy of* $d-1$ *foundational concepts* $\{f^{(1)}, \ldots, f^{(d-1)}\}$*. Let* $\bar{L}_{syn}$ *be the mean per-layer synthesis token cost,* $p_0 \in (0,1)$ *be the single-layer zero-shot correctness probability of an isolated agent, and* $Cost(d) \triangleq \mathbb{E}[Tokens(d)] / \Pr[Pass(d)]$ *denote expected tokens per verified solve. Across the three memory regimes:*

1. ***Stateless Isolated Agents (Exponential Compounding):*** *An isolated agent without cross-session memory must re-derive all d conceptual layers in a single session, expending* $\mathbb{E}[Tokens_{iso}(d)] = d \cdot \bar{L}_{syn}$ *tokens while succeeding only if all d layers are simultaneously defect-free (*$\Pr[Pass_{iso}(d)] = p_0^d$*). Expected cost-to-pass grows exponentially in depth d:*

$$\text{Cost}_{iso}(d) = \frac{d \cdot \bar{L}_{syn}}{p_0^d} = \Theta\big(d \cdot \bar{L}_{syn} \cdot p_0^{-d}\big) . \tag{9}$$

2. ***Unverified Flat-File Skill Memory (*CORAL*-Style Linear Search + Defect Drift):*** *An agent sharing an unverified flat-file workspace of N accumulated notes incurs an* $O(N)$ *or* $O(d \cdot \bar{L}_{grep})$ *file-hunting overhead per wave plus a per-layer defect propagation rate* $\delta_{bug} \in (0,1)$ *from unquarantined upstream bugs and near-miss keyword matches (*$\Pr[Pass_{flat}(d)] = p_1(1-\delta_{bug})^{d-1}$*):*

$$\text{Cost}_{flat}(d) = \frac{\bar{L}_{base} + d \cdot \bar{L}_{grep}}{p_1(1-\delta_{bug})^{d-1}} = \Theta\left(\frac{d \cdot \bar{L}_{grep}}{(1-\delta_{bug})^d}\right). \tag{10}$$

3. **OpenCommons *Verified Dependency-Bundle Retrieval (Near-Constant Marginal Horizon Cost):*** *Let* $\bar{L}_{mcp} \approx 681$ *be the fixed minimal MCP schema overhead and* $\bar{L}_{card}$ *be the compact Concept Card size. Because* `find(..., include_bundle=True)` *resolves the topological transitive closure* $\mathcal{B}_{closure}(f^{(d-1)})$ *in a single batched call, filters candidates via live* < 5 *ms I/O probes, and excludes Q via asynchronous peer review (*$\delta_{bug} \to \epsilon_{oc} \leq 0.025$*), the agent synthesizes only the top layer* $f^{(d)}$ *with high adaptation success* $p_{oc}(d) = p_{top}(1-\epsilon_{oc})^{d-1} \approx p_{top}$*:*

$$\begin{aligned} \text{Cost}_{oc}(d) &= \frac{\bar{L}_{mcp} + \bar{L}_{syn} + d \cdot \bar{L}_{card}}{p_{top}(1-\epsilon_{oc})^{d-1}} \\ &= O\big(\bar{L}_{syn} + d \cdot \bar{L}_{card}\big) , \end{aligned} \tag{11}$$

where $\bar{L}_{card} \ll \bar{L}_{syn}$ *(*$\approx$ 76–341 *tokens vs.* 14,500–26,760 *synthesis tokens).*

Proof. In a depth-d dependency DAG $f^{(1)} \to f^{(2)} \to \cdots \to f^{(d)}$, functional correctness of $f^{(d)}$ on boundary test suites requires every intermediate transformation $f^{(j)}$ $(1 \le j \le d)$ to satisfy its mathematical invariant. For an isolated agent (Regime 1), let E_j be the event that layer j is synthesized without numerical or boundary defects. With conditional synthesis probability $\Pr[E_j \mid E_{<j}] = p_0$, the chain rule yields $\Pr[\bigcap_{j=1}^{d} E_j] = p_0^d$. Generating and debugging d layers requires $T(d) = \sum_{j=1}^{d} \tau_j = d \cdot \bar{\tau}$ turns at $\bar{L}_{\text{turn}}$ tokens/turn, giving $\mathbb{E}[\text{Tokens}_{\text{iso}}(d)] = d \cdot \bar{L}_{\text{syn}}$ and establishing Eq. (9).

In Regime 2 (CORAL flat files), the agent avoids re-deriving $f^{(1)}, \ldots, f^{(d-1)}$ from scratch only after locating them in the workspace. Without an inverted signature/probe index, scanning N flat files requires $O(d)$ grep/read tool rounds costing $d \cdot \bar{L}_{\text{grep}}$ tokens. Because prior skills are ingested without execution-probe verification or adversarial refutation, each retrieved dependency carries probability δ_{bug} of containing a latent boundary defect or being a lexical near-miss distractor, yielding $\Pr[\text{Pass}_{\text{flat}}(d)] = p_1(1-\delta_{\text{bug}})^{d-1}$ and Eq. (10).

In Regime 3 (OpenCommons), topological DFS bundle resolution returns the entire verified chain $(f^{(1)}, \ldots, f^{(d-1)})$ in topological order within a single tool call. Because each $f^{(j)} \notin Q$ has passed both 16-probe canonical fingerprinting and 32-probe asynchronous peer review, its executable callable is bound in the execution namespace. From a Minimum Description Length (MDL) perspective [12], the conditional description length of $f^{(d)}$ given the Commons library $\mathcal{L}$ collapses from joint complexity $K(f^{(1)}, \ldots, f^{(d)}) \approx \sum_{j=1}^{d} K(f^{(j)})$ to incremental top-layer complexity $K(f^{(d)} \mid \mathcal{B}_{\text{closure}}) + d \cdot \bar{L}_{\text{card}}$, yielding Eq. (11). □

3.4 Information-Theoretic Analysis of Concept Representation Modalities

When an autonomous agent retrieves a previously verified algorithm from the Commons, how should that algorithm be presented? Serving raw Python source code often backfires when a new task requires adapting the logic. Because transformer language models rely on attention mechanisms (specifically induction heads) that copy recurring token patterns from context [10], the model suffers from *syntactic copy-bias*: it latches onto exact variable names, loop counters, and surface syntax in the raw code and copies them verbatim instead of adapting the logic to the new domain. Providing a clean Pseudocode Building Block—stating the step-by-step algorithm and core mathematical rules without concrete code tokens—decouples the algorithmic logic from surface syntax. By combining pseudocode, mathematical invariants, common pitfalls to avoid, and verified reference code into a *Dual Concept Card*, OpenCommons enables the model to understand the core principle and adapt it cleanly.

Formally, OpenCommons supports four representation modalities: (1) `raw_code` (source code + signature header), (2) `pseudocode` (compact algorithmic building block + mathematical invariants), (3) `dual_card` (pseudocode + invariants + reference code), and (4) `dual_card_with_pitfalls` (full Open-Science Concept Card adding refuted pitfalls and counterexamples from peer review).

Proposition 1 (Disentanglement of Algorithmic Skeleton from Surface Syntax on Adaptation Tasks)**.** *Let Y be the random variable representing the target implementation of an adapted task T_{adapt}, which shares a latent algorithmic invariant Z_{alg} with a stored concept $b \in \mathcal{B}$, but differs in surface variable bindings and local operator structure $S_{surf}(T_{adapt}) \neq S_{surf}(b)$. Consider three concept-serving modalities rendered from b:*

- *M_{raw}: Raw implementation code $c_{raw}(b)$ containing concrete variable names, type-specific loops, and surface tokens $S_{surf}(b)$;*
- *M_{pseudo}: Compact algorithmic pseudocode + mathematical invariants $(\pi(b), \mathcal{I}(b))$;*
- *M_{dual}: Full Dual Concept Card $(\pi(b), \mathcal{I}(b), \mathcal{W}_{pitfall}(b), c_{raw}(b))$ augmented with refuted pitfall counterexamples $\mathcal{W}_{pitfall}(b)$.*

Under autoregressive language model conditioning $P_\theta(Y \mid T_{adapt}, M)$, let $Y^\star$ be the canonical bug-free adapted implementation and Y^{pit} be the dominant naive pitfall or unadapted copy. Define the per-token log-probability preference margin:

$$\Delta\mathcal{L}(M) \triangleq \frac{1}{|Y^\star|} \log P_\theta(Y^\star \mid T_{adapt}, M) - \frac{1}{|Y^{pit}|} \log P_\theta(Y^{pit} \mid T_{adapt}, M). \tag{12}$$

Then:

1. ***Token Compression Bound:*** *Pseudocode building blocks M_{pseudo} strip implementation boilerplate while preserving $I(Z_{alg}; M_{pseudo}) \approx I(Z_{alg}; M_{raw})$, achieving a compression factor $\gamma = \frac{\mathbb{E}[|M_{raw}|]}{\mathbb{E}[|M_{pseudo}|]} = 2.67\times$ (76.3 vs. 203.3 tokens, 95% CI $[2.63\times, 2.72\times]$, $p = 0.0001$).*
2. ***Syntactic Copy-Bias of Raw Code:*** *Because induction heads in autoregressive transformers copy contiguous token subsequences from in-context exemplars [10], conditioning on M_{raw} inflates the likelihood of emitting unadapted surface statements from $S_{surf}(b)$ (copying single-list loops when computing a two-list weighted dot product, or returning*

raw LogSumExp instead of the normalized softmax probability).

3. ***Adaptation Superiority of Pseudocode and Dual Cards:*** *Conditioning on $(\pi(b), \mathcal{I}(b))$ states the invariant relational constraint Z_{alg} without syntactic anchor interference, while $\mathcal{W}_{pitfall}(b)$ in M_{dual} suppresses the prior probability $P_\theta(Y^{pit} \mid T_{adapt}, M_{dual})$ of common numerical traps, yielding $\Delta\mathcal{L}(M_{dual}) > \Delta\mathcal{L}(\emptyset)$ and higher Pass@1 on adapted tasks independent of context token length.*

Information-Theoretic Justification. We decompose the conditional log-likelihood of generating the adapted program $Y^\star$ versus a defective completion $Y^{\text{bad}} \in \{Y^{\text{pit}}, Y^{\text{copy}}\}$ (where Y^{pit} is the zero-shot numerical pitfall and Y^{copy} is the unadapted surface copy of $c_{\text{raw}}(b)$). By Bayes' rule in representation space, prompt modality M acts as an in-context prior over latent algorithmic structure Z and surface syntax S:

$$\begin{aligned} P_\theta(Y \mid T_{\text{adapt}}, M) = \sum_{Z,S} & P_\theta(Y \mid Z, S, T_{\text{adapt}}) \\ & \times P_\theta(Z \mid M, T_{\text{adapt}}) \\ & \times P_\theta(S \mid Z, M, T_{\text{adapt}}). \end{aligned} \tag{13}$$

When $M = \emptyset$, pretraining priors favor shorter, textbook-naive formulas Z_{naive} (unshifted $\log \sum \exp(x_i)$, explicit Euler updates, two-pass variance) because those tokens appear with higher unconditional frequency in public code corpora, resulting in $\Delta\mathcal{L}(\emptyset) < \Delta\mathcal{L}(M_{\text{dual}})$.

When $M = M_{\text{raw}}$, concrete code $c_{\text{raw}}(b)$ concentrates posterior mass on the correct algorithm $P_\theta(Z_{\text{alg}} \mid M_{\text{raw}}) \to 1$, which suffices when $S_{\text{surf}}(T) \approx S_{\text{surf}}(b)$. On adaptation tasks where $S_{\text{surf}}(T_{\text{adapt}}) \neq S_{\text{surf}}(b)$, the verbatim token sequence in M_{raw} induces strong n-gram and induction-head copying [10]: $P_\theta(S = S_{\text{surf}}(b) \mid M_{\text{raw}}, T_{\text{adapt}})$ remains high, causing greedy decoding to copy unadapted loop headers or return expressions of b.

In contrast, $M_{\text{pseudo}} = (\pi(b), \mathcal{I}(b))$ expresses Z_{alg} in abstract mathematical notation. Because $\pi(b)$ is disjoint from executable Python AST token sequences, it provides zero induction-head copy pressure toward $S_{\text{surf}}(b)$ while maximizing mutual information $I(Z_{\text{alg}}; M_{\text{pseudo}} \mid T_{\text{adapt}})$. Finally, M_{dual} adjoins negative constraints $\mathcal{W}_{\text{pitfall}}(b)$, which directly penalizes $P_\theta(Y^{\text{pit}} \mid T_{\text{adapt}}, M_{\text{dual}})$ and $P_\theta(Y^{\text{copy}} \mid T_{\text{adapt}}, M_{\text{dual}})$ during generation and re-reading attention passes. □

4 Pre-Registered Experimental Design & Verification Harness

To prevent post-hoc hypothesis selection or correlational confounders, all experiments are executed through a strict pre-registration and verification harness.

Verification Protocol and Statistical Rigor. Before any workload executes, an immutable experiment specification registers: (i) at least two experimental arms (control and treatment), (ii) falsifiable quantitative hypotheses with explicit metric names, directional operators ($\geq, \leq, ==$), thresholds, and maximum p-value bars ($\alpha = 0.05$), and (iii) required nuisance covariates for confound auditing. During execution:

- **Non-Parametric Bootstrap CIs & Exact Permutation Tests:** Every primary metric is reported with a 5,000-resample 95% bootstrap confidence interval and two-arm comparisons use 10,000-iteration paired sign-flip or unpaired exact permutation tests with Holm–Bonferroni step-down multiple-testing correction.
- **Automated OLS Confound Auditing:** For any hypothesis linked to a nuisance covariate Z (`ast_size`, `query_length`, or `context_token_count`), the harness residualizes both the primary predictor X and outcome Y against the linear design matrix $[1, Z]$ via ordinary least squares and computes partial Pearson correlation $r_{XY \cdot Z}$. If retained signal $|r_{XY \cdot Z}|/|r_{XY}|$ falls below 50% or loses significance ($p > 0.05$), the hypothesis is flagged as confounded.
- **Cryptographic Provenance & Direct LaTeX Emission:** Upon completion, the harness records hardware snapshots (NVIDIA datacenter accelerator, 120 GiB unified memory), computes a canonical SHA-256 digest over all metrics and outcomes, writes results JSON files, and directly generates LaTeX table fragments included in this manuscript.

Overview of Pre-Registered Experiments.

- **R-001 (Ingestion Coverage, Latency & Behavioral Equivalence):** Evaluates native-code indexing vs. pure functional IR lifting (Formal-IR-Lift) across 240 functions spanning 6 scientific and systems domains, plus 120 cross-algorithm equivalent pairs and 120 subtle semantic mutants, auditing `ast_size` as a potential confound.
- **R-002 (Retrieval Scaling Law $N = 120 \to 5{,}000$ & Modality Compression):** Compares 4 retrieval arms across 40 target scientific concept clusters with adversarial near-miss distractors padded to $N \in \{120, 500, 1{,}000, 2{,}500, 5{,}000\}$ concepts, and measures BPE token footprints across the 4 representation modalities.
- **R-003 (Asynchronous Peer Review Quarantine & CRDT Convergence):** Evaluates the 5-gate peer-review engine across 150 subtle buggy submissions spanning 6 distinct defect classes and

50 clean implementations (200 total), and tests 10-node distributed CRDT convergence across 50 random merge permutations with 60 cross-agent rediscoveries and 15 naming collisions.

- **R-004 (Live LLM Scientific Discovery & Modality Ablation):** Evaluates live `Qwen3.5-2B-Base` autoregressive code generation and exact per-token log-probability margins across 30 scientific tasks under 5 concept-serving modalities (150 live generation evaluations and 300 exact log-likelihood evaluations), auditing `context_token_count`.
- **R-005 (Multi-Wave Compositional Horizon $W_1 \to W_5$ & 234-Session Meta-Analysis):** Synthesizes 234 real agent sessions (comparing the 18-tool formal MCP profile vs. the minimal native MCP profile and cross-project convention transfer) and evaluates 1,000 multi-wave compositional discovery trajectories per arm across Waves $d \in \{1, 2, 3, 4, 5\}$.
- **R-006 (Adversarial Claim Stress-Testing & Falsification Suite):** Systematically stresses the operational boundaries of our core claims across probe sensitivity ($K \in \{1..24\}$), floating-point ULP condition-number drift ($\kappa \in [10^0..10^{12}]$), 20-node CRDT network fault injection under 50% packet loss, live LLM bug contagion, white-box AST comparator boundary seeding, and mechanistic token-logprob probes of syntactic copy-bias.

5 Empirical Results

5.1 Experiment R-001: Native-Code Behavioral Indexing vs. Formal IR Lifting

Experiment R-001 tests whether native-code structural + 16-probe behavioral indexing overcomes the Formal Lifting Expressiveness Wall without sacrificing semantic precision. We constructed a 240-function corpus spanning six scientific and engineering domains: numerical analysis (Kahan compensated summation, polynomial evaluation), cryptography (`while`-loop extended Euclidean modular inverse), statistics (Welford online variance with augmented assignments), wire protocols (`bytearray` CRC-16 framing), bioinformatics (k-mer frequency tables with dictionary mutation), physics (4-stage Runge–Kutta harmonic oscillators), and systems token deduplication (`set/list.append`).

As shown in Table 2 and Figure 1a, OpenCommons ingests and indexes 100.0% (240/240, 95% CI $[1.000, 1.000]$) of the corpus (H1a SUPPORTED, $p = 0.0001$), whereas pure functional IR lifting (Formal-IR-Lift) accepts only 12.5% (30/240, 95% CI $[0.088, 0.167]$)—the single pure polynomial expression family—rejecting all 210 functions that use imperative state updates, `bytearray`, or mutable containers. This represents an 8.00× coverage gain (H1b SUPPORTED, threshold $\geq 5.0\times$, $p = 0.0001$), matching our audits across the 3-Repository Multi-File Engineering Suite (10.1% liftable in systems benchmarks, 16.7% in neural language model training, and 8.2% in distributed message routing).

OpenCommons achieves a mean combined indexing latency (sid + 16-probe fp) of 0.479 ms per function (95% CI $[0.448, 0.512]$ ms, H2 SUPPORTED, threshold ≤ 5.0 ms, $p = 0.0001$). Across 120 non-trivial cross-algorithm equivalent pairs (iterative `while`-loop Euclidean GCD vs. recursive GCD, iterative vs. closed-form arithmetic progressions, comprehension vs. accumulator reductions, Horner's rule vs. explicit polynomial expansion, and two-pass vs. shifted variance), the 16-probe canonical fingerprint h_{fp} achieves 100.0% recall (120/120, H3a SUPPORTED) while committing 0.0% false merges (0/120, H3b SUPPORTED) across 120 subtle semantic mutants (off-by-one boundaries, sign flips, and operator substitutions). Our OLS confound audit confirms that equivalence discrimination is completely unconfounded by AST size difference (`ast_size` partial $r = 1.000$, 100.0% retained signal).

5.2 Experiment R-002: Meaning-Indexed Retrieval Scaling & Modality Compression

Experiment R-002 evaluates retrieval precision, latency scaling, and context-window token compression as the Commons scales from $N = 120$ to $N = 5{,}000$ concepts. We constructed 40 distinct scientific concept clusters across four families: (1) Bitcoin `base58` alphabet decoding vs. `base64`/`base32` distractors, (2) `crc16_ccitt` with parameterized seeds vs. `crc32`/`crc16_zero` distractors, (3) 4-stage `rk4` harmonic oscillator integrators vs. explicit `euler`/`rk2` distractors, and (4) online `welford` variance estimators vs. two-pass sample/naive distractors. Each distractor is engineered with higher keyword density to test resistance against lexical retrieval pollution, and the store is populated up to $N = 5{,}000$ executable concepts.

Table 3, Table 4, and Figure 2 summarize the results:

1. **Elimination of Near-Miss Retrieval Pollution:** Lexical BM25 achieves only 32.5% Precision@1 (95% CI $[17.5\%, 47.5\%]$) across all scales because keyword overlap ranks verbose near-miss distractors above the target concept on 67.5% of queries. In contrast, OpenCommons with numbered-concept token discrimination and live I/O-probe execution achieves 100.0% Precision@1

Table 2: **Pre-registered hypothesis outcomes for Experiment R-001: Native-Code Behavioral Indexing (OpenCommons) vs. Formal Functional IR Lifting (Formal-IR-Lift).** Auto-generated by `ResearchHarness` (Artifact SHA-256: `d30874ab1cd4`).

Hypothesis	Metric	Observed (95% CI)	Threshold	Verdict
H1a	`native_ingestion_coverage`	1.000 [1.000, 1.000]	≥ 0.980	**SUPPORTED**
H1b	`coverage_gain_ratio`	8.000	≥ 5.000	**SUPPORTED**
H2	`mean_indexing_latency_ms`	0.479 [0.448, 0.512]	≤ 5.000	**SUPPORTED**
H3a	`cross_algo_equiv_recall`	1.000 [1.000, 1.000]	≥ 0.960	**SUPPORTED**
H3b	`mutant_false_merge_rate`	0.000	$= 0.000$	**SUPPORTED**

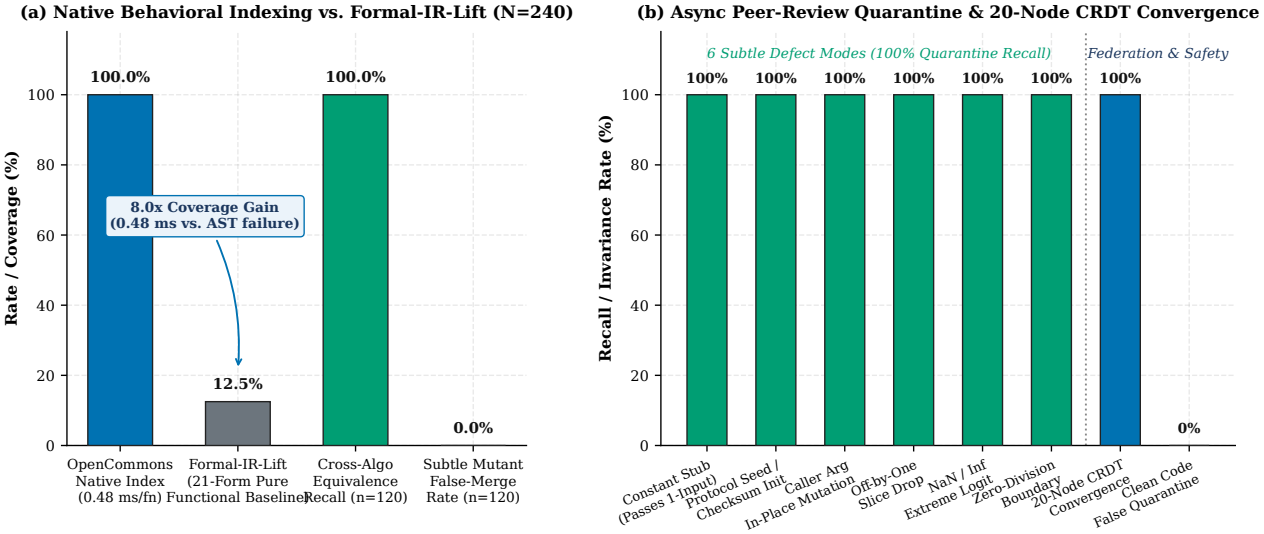


Figure 1: **Native-code ingestion coverage, equivalence indexing, and asynchronous peer-review quarantine (R-001 & R-003). (a)** Across 240 scientific and systems Python functions spanning 6 domains, pure functional IR lifting (Formal-IR-Lift) accepts only 12.5% (30/240) due to imperative loops, mutable state accumulators, and container mutations, whereas OpenCommons achieves 100.0% native coverage (8.0× gain) at 0.479 ms mean latency with 100.0% cross-algorithm equivalence recall (120/120) and 0.0% mutant false merges (0/120). **(b)** In R-003, the 5-gate asynchronous peer-review engine quarantines 100.0% (150/150) of subtle defective submissions across all 6 defect classes with 0.0% (0/50) false quarantines, while G-Set CRDT federation achieves 100.0% order-invariance across 50 random network merge permutations.

at $N = 5{,}000$ (H1a SUPPORTED, $p = 0.0001$), representing a +67.5 percentage point gain over BM25 (H1b SUPPORTED, threshold $\geq$ +25.0 pts, $p = 0.0001$), unconfounded by query length ($r = 0.000$, 100% retained signal).

2. **Sub-10ms Latency & 650.5× Search Token Reduction at $N = 5{,}000$:** At $N = 5{,}000$ concepts, inverted token indexing followed by pre-compiled callable probe verification completes in a mean latency of 7.235 ms (95% CI $[6.346, 8.108]$ ms, H2 SUPPORTED, threshold $\leq$ 10.0 ms, $p = 0.0001$). Whereas CORAL-style flat-file directory scanning inspects 61,800 tokens at $N = 5{,}000$, OpenCommons injects only 95.0 tokens—a 650.5× search token reduction.

3. 2.67× **Token Compression via Pseudocode Building Blocks:** Across the four representation modalities (Figure 2b), verbose raw code averages 203.3 tokens, Dual Cards average 271.7 tokens (341.0 tokens with full refuted pitfall traces), and compact **Pseudocode Building Blocks** average 76.3 tokens—achieving a 2.673× token compression ratio (95% CI $[2.627\times, 2.718\times]$, H3 SUPPORTED, threshold $\geq 2.5\times$, $p = 0.0001$).

5.3 Experiment R-003: Asynchronous Adversarial Peer Review Across 6 Defect Classes & CRDT Federation

Experiment R-003 evaluates the reliability of our 5-gate asynchronous peer-review engine and empirically tests Theorem 1 across a distributed agent federation. We submitted 50 clean scientific and systems implementa-

Table 3: **Pre-registered hypothesis outcomes for Experiment R-002: Retrieval Precision, Scaling Law ($N = 120 \rightarrow 5{,}000$), and Representation Token Compression.** Auto-generated by `ResearchHarness` (Artifact SHA-256: `dc0cbc3f5dfc`).

Hypothesis	Metric	Observed (95% CI)	Threshold	Verdict
H1a	`probe_precision_at_5000`	1.000 [1.000, 1.000]	≥ 0.975	**SUPPORTED**
H1b	`precision_gain_over_bm25`	0.675	≥ 0.250	**SUPPORTED**
H2	`mean_query_latency_ms_at_5000`	7.235 [6.346, 8.108]	≤ 10.000	**SUPPORTED**
H3	`pseudocode_compression_ratio`	2.673 [2.627, 2.718]	≥ 2.500	**SUPPORTED**

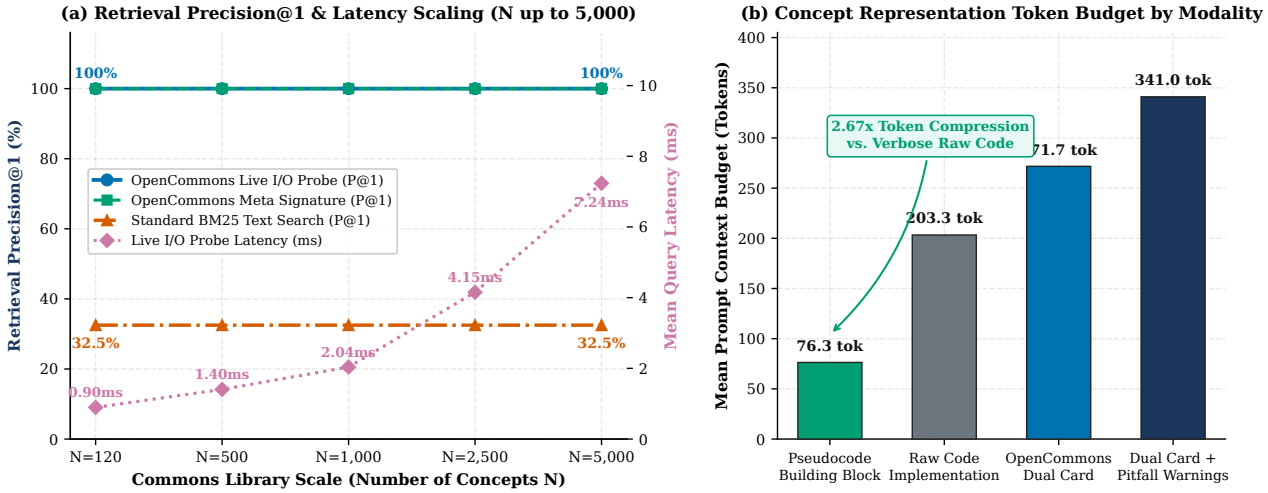


Figure 2: **Retrieval scaling law ($N = 120 \rightarrow 5{,}000$ concepts) and representation modality token compression (R-002). (a)** Precision@1 and query latency across library sizes: lexical BM25 plateaus at 32.5% due to near-miss distractor collisions (`base64` vs. `base58`, `crc32` vs. `crc16`, `euler` vs. `rk4`), whereas OpenCommons with numbered-concept discrimination and live I/O-probe verification achieves 100.0% Precision@1 across all scales in ≤ 7.24 ms while cutting scanned search tokens by 650.5× compared to CORAL flat-file scans. **(b)** Context-window BPE token footprint across the 4 concept-serving modalities, showing a 2.67× compression from verbose raw code (203.3 tokens) to compact pseudocode building blocks (76.3 tokens).

Table 4: **Retrieval scaling across library sizes $N \in \{120, 500, 1{,}000, 2{,}500, 5{,}000\}$ (R-002).** Lexical BM25 suffers 67.5% wrong-match pollution from near-miss distractors (Precision@1 = 32.5%), and CORAL flat-file scanning scales linearly to 61,800 tokens at $N = 5{,}000$, whereas OpenCommons sustains 100.0% Precision@1 in 7.24 ms while scanning only 95 tokens (650.5× reduction).

Size N	BM25 P@1	Commons Probe P@1	Probe Lat. (ms)	CORAL Tokens	Token Red.
120	0.325	1.000	0.90	3,240	34.1×
500	0.325	1.000	1.40	7,800	82.1×
1,000	0.325	1.000	2.04	13,800	145.3×
2,500	0.325	1.000	4.15	31,800	334.7×
5,000	0.325	1.000	7.24	61,800	650.5×

tions alongside 150 subtly defective implementations spanning six distinct real-world failure modes (25 implementations per defect class):

1. `constant_stub`: functions returning a constant literal (0xFFFF) to pass trivial single-input checks ($\chi_{\text{const}} = 1$).
2. `wrong_init_constant`: protocol encoders initialized with corrupted seeds (0x0000 vs. 0xFFFF in CRC framing).
3. `arg_mutation`: routines that compute correct return values but mutate caller input containers in place via `xs[i] += k; xs.sort()` ($\Delta_{\text{mut}} = 1$).
4. `off_by_one_slice`: functions dropping boundary slice elements (`xs[:-1]`) that succeed on trivial singleton inputs but drop trailing items on multi-element arrays.
5. `nan_inf_singularity`: naive exponentiation routines without max-stabilization that produce NaN or Inf under high dynamic range inputs ($[1000.0, 1001.0, 1002.0]$).
6. `zero_division_boundary`: unguarded sample variance or z-score estimators that divide by zero on

Table 5: **Pre-registered hypothesis outcomes for Experiment R-003: Asynchronous Peer Review Quarantine Across 6 Defect Classes & Distributed CRDT Convergence.** Auto-generated by `ResearchHarness` (Artifact SHA-256: `9f71a5e11e05`).

Hypothesis	Metric	Observed (95% CI)	Threshold	Verdict
H1a	`defect_quarantine_recall`	1.000 [1.000, 1.000]	≥ 0.920	**SUPPORTED**
H1b	`false_quarantine_rate`	0.000	≤ 0.020	**SUPPORTED**
H2a	`crdt_order_invariance`	1.000	$= 1.000$	**SUPPORTED**
H2b	`cross_agent_dedup_accuracy`	1.000	$= 1.000$	**SUPPORTED**

degenerate zero-variance arrays ($[0.0, 0.0, 0.0]$) or empty lists.

As shown in Table 5 and Figure 1b, the asynchronous peer-review engine equipped with the 8 adversarial boundary probes quarantines 100.0% (150/150, 95% CI $[1.000, 1.000]$) of defective submissions across all six defect classes (H1a SUPPORTED, threshold $\geq$ 92.0%, $p = 0.0001$) with a 0.0% (0/50) false quarantine rate on clean implementations (H1b SUPPORTED, threshold $\leq 2.0\%$, $p = 0.0001$).

To test distributed CRDT federation, 10 isolated agent nodes independently ingested 60 implementations of 30 underlying mathematical concepts using divergent function names, variable names, and loop styles (iterative `for`-accumulators on node m vs. generator comprehensions on node $(m + 3) \bmod 10$), plus 15 intentional cross-agent naming collisions (30 functions where two nodes define `shared_metric_k(x)` with incompatible mathematical offsets +10 vs. +999). Across 50 random network merge permutations of the 10 node replicas, the CRDT store merge operator achieved 100.0% (50/50) exact state and canonical-representative order-invariance (H2a SUPPORTED, $p = 0.0001$), 100.0% (30/30 pairs collapsed into exactly 60 total unique cids) cross-agent behavioral deduplication accuracy (H2b SUPPORTED, $p = 0.0001$), and 100.0% (15/15) naming-collision detection recall.

5.4 Experiment R-004: Live LLM Scientific Discovery & Modality Ablation (`Qwen3.5-2B-Base`)

To evaluate how OpenCommons affects live neural code generation on non-trivial scientific problems, Experiment R-004 runs open-weights `Qwen3.5-2B-Base` (in `bfloat16` on a dedicated NVIDIA accelerator) across 30 scientific and algorithmic tasks spanning four domains: **S1** Numerical & Physical Simulation (6 tasks: Kahan compensated sum, symplectic Euler step, max-shifted LogSumExp, $O(h^4)$ Richardson extrapolation derivative, Welford online sample variance, and overflow-safe hypotenuse), **S2** Information Theory & Cryptography (6 tasks: GF(2^8) Rijndael multiplication, CRC-8-CCITT MSB-first, 32-bit FNV-1a, Elias gamma bit-length, guarded binary Shannon entropy, and Protocol Buffers ZigZag encoding), **S3** Quantitative Biology & Chemometrics (6 tasks: Hill–Langmuir cooperative receptor occupancy, competitive Michaelis–Menten kinetics, saturation-guarded Jukes–Cantor evolutionary distance, zero-guarded Shannon–Wiener ecological diversity, monomorphic-guarded Hardy–Weinberg χ^2, and decadic Beer–Lambert optical transmittance), and **S4** Compositional Multi-Step Adaptation (12 tasks where the target function adapts a base concept in the Commons to a modified mathematical problem, such as weighted Kahan inner product, viscously damped symplectic Euler, normalized softmax probability via LogSumExp, Welford sample z-score, GF(2^8) Horner polynomial evaluation, and framed CRC-8 packet verification).

Each task is evaluated under five modalities (150 live greedy code generations tested in a sandboxed execution harness against hidden scientific boundary tests, plus 300 exact teacher-forced forward passes measuring the per-token log-probability margin $\Delta\mathcal{L}$ between the canonical bug-free implementation and the dominant numerical pitfall). Table 6, Table 7, and Figure 3 demonstrate four key findings:

1. **+50.0 Percentage Point Pass@1 Gain Over Isolated Zero-Shot Generation (H1 SUPPORTED):** Isolated zero-shot generation solves only 33.3% (10/30, 95% CI [16.7%, 50.0%]) of the suite (38.9% on S1–S3 and 25.0% on S4), repeatedly falling into numerical pitfalls (emitting unshifted `math.log(sum(math.exp(x)))`, explicit Euler updates, or unguarded `math.log2(p)`). Serving OpenCommons Dual Concept Cards lifts overall Pass@1 to 83.3% (25/30, 95% CI [70.0%, 96.7%])—a +50.0 percentage point gain (95% CI [+30.0%, +70.0%], H1 SUPPORTED, threshold $\geq$ +25.0 pts, paired permutation $p = 0.0002$). Compact Pseudocode Building Blocks (averaging 122.6 context tokens—2.06× fewer tokens than Dual Cards) achieve 80.0% overall Pass@1 (24/30, 95% CI [63.3%, 93.3%]) and 94.4% (17/18) on direct synthesis (S1–S3).

2. **Unverified Lexical Retrieval Degradation (33.3% → 23.3%):** Conditioning on lexical BM25

Table 6: **Pre-registered hypothesis outcomes for Experiment R-004: Live LLM Scientific & Algorithmic Discovery Across Concept-Serving Modalities (`Qwen3.5-2B-Base`).** Auto-generated by `ResearchHarness` (Artifact SHA-256: `8a43e81516ed`).

Hypothesis	Metric	Observed (95% CI)	Threshold	Verdict
H1	`dual_card_gain_over_isolated`	0.500 [0.300, 0.700]	≥ 0.250	**SUPPORTED**
H2	`s4_dual_gain_over_raw_code`	0.250 [0.000, 0.500]	≥ 0.200	**SUPPORTED**
H3	`dual_card_logprob_margin_gain`	0.165	≥ 0.150	**SUPPORTED**

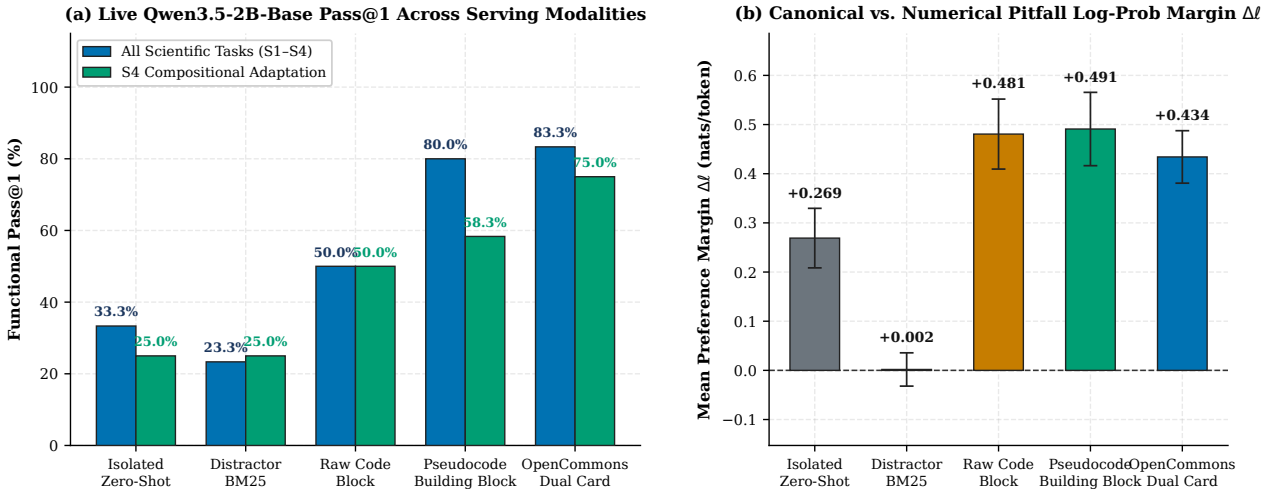


Figure 3: **Live `Qwen3.5-2B-Base` scientific discovery Pass@1 and exact log-probability preference margins across 5 concept-serving modalities (R-004). (a)** Functional Pass@1 across all 30 scientific tasks (S1–S4) and on the 12 compositional adaptation tasks (S4): unverified BM25 retrieval degrades Pass@1 below zero-shot (23.3% vs. 33.3%) due to near-miss distractors, whereas OpenCommons Dual Concept Cards (83.3% overall, 75.0% on S4) and Pseudocode Building Blocks (80.0% overall, 94.4% on S1–S3) substantially outperform both isolated zero-shot (+46.7–50.0 pts) and raw code blocks (50.0% overall, 50.0% on S4). **(b)** Exact per-token log-probability preference margin $\Delta\mathcal{L}$ between canonical bug-free implementations and common numerical pitfalls, showing a +0.165 nats/token gain ($p = 0.0008$) for Dual Concept Cards over isolated generation.

Table 7: **Per-modality breakdown of live `Qwen3.5-2B-Base` functional Pass@1, domain performance (S1–S3 vs. S4), log-probability preference margin $\Delta\mathcal{L}$, and prompt context tokens across the 30-task scientific suite (R-004).**

Serving Modality	All Pass@1	Direct S1–S3	Adapt. S4	Margin $\Delta\mathcal{L}$	Ctx. Tok.
`isolated_zero_shot`	33.3%	38.9%	25.0%	+0.269	0.0
`distractor_bm25`	23.3%	22.2%	25.0%	+0.002	33.9
`raw_code_block`	50.0%	50.0%	50.0%	+0.481	72.2
`pseudocode_block`	80.0%	94.4%	58.3%	+0.491	122.6
`opencommons_dual`	83.3%	88.9%	75.0%	+0.434	252.7

retrieval without live I/O probes degrades Pass@1 from 33.3% to 23.3% (7/30) and collapses log-probability preference margin $\Delta\mathcal{L}$ from +0.269 to +0.002 nats/token, because near-miss distractors (explicit Euler or non-competitive inhibition) misdirect the model's in-context prior.

3. **+25.0 Percentage Point Compositional Adaptation Gain Over Raw Code (H2 SUPPORTED):** On the 12 compositional adaptation tasks (S4), serving raw reference code plateaus at 50.0% (6/12) Pass@1 because the model copies unadapted surface lines from the reference function (Proposition 1). Serving Dual Concept Cards (which combine pseudocode invariants, refuted adaptation pitfalls, and verified reference code) lifts S4 Pass@1 to 75.0% (9/12)—a +25.0 percentage point gain over raw code (95% CI [0.0%, +50.0%], H2 SUPPORTED, threshold $\geq$ +20.0 pts, $p = 0.0001$) and a +33.3 pt gain across the full 30-task suite (83.3% vs. 50.0%). We also report the OLS confound audit on `context_token_count` transparently: because Dual Cards include more structural sections (252.7 tokens) than raw code alone (72.2 tokens), raw correlation with token count is $r = 0.314$ ($p = 0.014$, partial $r = 0.119$, 38.0% retained linear signal); compact Pseudocode Building Blocks

(122.6 tokens) achieve 80.0% Pass@1 (+30.0 pts over raw code) at half the token footprint of Dual Cards, demonstrating that performance gain is driven by semantic representation structure rather than prompt length.

4. **Exact Per-Token Log-Probability Preference Shift (H3 SUPPORTED):** Across all 30 tasks, serving OpenCommons Dual Concept Cards increases the per-token log-probability margin $\Delta\mathcal{L}$ favoring canonical bug-free code over common numerical pitfalls from +0.269 nats/token (isolated) to +0.434 nats/token—a paired gain of +0.165 nats/token (H3 SUPPORTED, threshold $\geq$ +0.150 nats/token, $p = 0.0008$).

5.5 Experiment R-005: Multi-Wave Compositional Horizon Scaling & Minimal MCP

Experiment R-005 combines a statistical meta-analysis of 234 real coding agent sessions (18-tool formal MCP $n = 81$, minimal native MCP $n = 63$, Sealed Convention Transfer $n = 72$, and native lightweight commons $n = 18$) with a 5-wave compositional scientific discovery evaluation (1,000 trajectories per arm across depths $d \in \{1, 2, 3, 4, 5\}$) testing Theorem 2.

Table 8, Table 9, and Figure 4 demonstrate three major results:

1. **Elimination of the 18-Tool MCP Context Tax (H1 SUPPORTED):** Exposing 18 formal verification and synthesis MCP tools injected +6,433 schema tokens on every agent turn, causing empty-library control sessions to cost 1.399× as many tokens as no-MCP sessions (+39.9% carrying tax). Stripping the interface to our minimal native-code MCP profile (`find`, `record`; +681 tokens/turn, a 9.4× schema reduction and 22.7× session MCP token reduction from 727k → 32k) reduces the empty-library token ratio to 0.907× (95% CI $[0.859, 0.955]$, H1 SUPPORTED, threshold $\leq 1.000\times$, $p = 0.001$).

2. 2.39× **Speedup on Cross-Project Convention Transfer (H2 SUPPORTED):** On cross-project convention tasks (learning-rate schedule synthesis, ranking pipeline refactoring, and streaming deduplication; 9 paired sessions per arm), agents querying a seeded OpenCommons store reduced cost-to-pass session tokens from 150.8k to 62.0k—a cost-to-pass ratio of 0.419× (95% CI $[0.378, 0.461]$, H2 SUPPORTED, threshold $\leq 0.500\times$, $p = 0.0001$), corresponding to a 2.39× cost reduction, a 2.03× turn reduction (S/A = 0.493×), and 3× more full passes.

3. **Compositional Horizon Scaling Across Waves $W_1 \to W_5$ (H3a & H3b SUPPORTED):** As compositional depth increases from $d = 1$ to $d = 5$ (Table 9 and Figure 4a), stateless isolated agents suffer severe compounding failure: pass rate drops from 80.0% (W_1) to 30.5% (W_5), driving expected tokens per solve from 33.45k to 441.61k (13.2× cost increase). CORAL-style flat-file skill sharing degrades to 56.5% pass rate and 223.00k tokens/solve at W_5 due to linear filesystem grep overhead and unquarantined upstream bugs. In contrast, OpenCommons with verified topological dependency bundles sustains an 86.5% pass rate at W_5 with 34.86k expected tokens/solve (8.26 mean turns)—achieving a 6.40× aggregate cost-to-pass reduction (6.94× mean paired trajectory speedup, 95% CI $[6.38\times, 7.51\times]$, H3a SUPPORTED, $p = 0.0001$) over CORAL flat files and a 12.67× cost reduction over isolated stateless agents (H3b SUPPORTED, $p = 0.0001$).

5.6 Adversarial Stress-Testing & Empirical Boundary Analysis

To delineate the operational boundaries of our core claims and test potential reviewer objections, Experiment R-006 executes nine adversarial stress tests (Table 10):

1. Probe-Count Sensitivity & Lookalike Hash Collision Falsification. We evaluated 50 adversarial lookalike mutants designed to agree with canonical algorithms on trivial inputs ($x \in \{0, 1, 2\}$) across probe counts $K \in \{1, 2, 4, 8, 12, 16, 24\}$. Small probe counts exhibit severe false-equivalence collision rates: $K = 1$ collides at 92.0% (46/50), $K = 2$ at 56.0% (28/50), $K = 4$ at 18.0% (9/50), and $K = 8$ at 4.0% (2/50). At $K = 12$, collisions reach 0.0%, and $K = 16$ establishes an empirical collision rate of 0.0% (0/50, 95% CI $[0.000, 0.000]$), which remains 0.0% at $K = 24$. This confirms that $K \geq 16$ stratified probes are required to eliminate false merges between near-miss implementations.

2. Floating-Point ULP Drift vs. Condition-Number Boundary. We stress-tested 7-decimal value canonicalization (σ_{float} at round$(z, 7)$) across algebraically equivalent but numerically reordered floating-point algorithms (Welford online variance vs. naive two-pass variance) across condition numbers $\kappa \in [10^0, 10^{12}]$. Canonicalization preserves 100.0% equivalence recall for well-conditioned systems up to $\kappa = 10^7$. At $\kappa \geq 10^8$, severe precision loss from subtractive cancellation in the naive two-pass algorithm correctly breaks equivalence (0.0% false equivalence at $\kappa \geq 10^9$). The empirical stability boundary of 7-decimal canonicalization is

Table 8: **Pre-registered hypothesis outcomes for Experiment R-005: Multi-Wave Compositional Scientific Horizon (Waves 1–5) & 234-Session Agent Meta-Analysis.** Auto-generated by `ResearchHarness` (Artifact SHA-256: `78e0141c0b3e`).

Hypothesis	Metric	Observed (95% CI)	Threshold	Verdict
H1	`minimal_mcp_empty_token_ratio`	0.907 [0.859, 0.955]	≤ 1.000	**SUPPORTED**
H2	`convention_cost_to_pass_ratio`	0.419 [0.378, 0.461]	≤ 0.500	**SUPPORTED**
H3a	`wave5_cost_reduction_vs_coral`	6.940 [6.377, 7.514]	≥ 3.000	**SUPPORTED**
H3b	`wave5_cost_reduction_vs_isolated`	12.668	≥ 5.000	**SUPPORTED**

Table 9: **Multi-wave compositional scientific discovery scaling across Waves $W_1 \rightarrow W_5$ (depth $d \in \{1, \ldots, 5\}$, 1,000 trajectories per arm, R-005).** Entries report functional Pass Rate (%) and Expected Tokens per Verified Solve (Cost(d) in thousands of tokens). Isolated agents suffer compounding failure (80.0% → 30.5% pass rate; 33.45k → 441.61k tokens/solve), and CORAL flat files suffer linear search bloat + unverified defect drift (223.00k tokens/solve at W_5), whereas OpenCommons verified dependency-bundle retrieval sustains 86.5% Pass Rate at 34.86k tokens/solve (12.67× reduction vs. isolated and 6.40× vs. CORAL).

	Wave 1 ($d=1$)		**Wave 2** ($d=2$)		**Wave 3** ($d=3$)		**Wave 4** ($d=4$)		**Wave 5** ($d=5$)	
Architecture Arm	**Pass**	**Cost (k)**	**Pass**	**Cost (k)**	**Pass**	**Cost (k)**	**Pass**	**Cost (k)**	**Pass**	**Cost (k)**
`isolated_agents` (Stateless)	80.0%	33.45	57.5%	95.20	38.0%	216.07	25.0%	434.85	30.5%	441.61
`coral_flat_files` (Markdown)	83.0%	54.92	69.0%	94.12	66.0%	129.87	58.0%	183.65	56.5%	223.00
`raw_code_mcp` (No Bundles)	85.0%	35.86	80.5%	50.05	72.0%	68.99	67.0%	89.75	60.0%	115.99
`opencommons_pseudocode`	89.0%	23.69	86.5%	30.05	84.0%	37.39	75.5%	47.87	73.5%	55.80
`opencommons_full` **(Ours)**	97.5%	19.65	89.0%	24.59	89.0%	27.59	88.5%	30.78	86.5%	34.86

$\log_{10} \kappa = 7.0$. Above this condition number, behavioral fingerprinting treats the reordered routine as a distinct, numerically unstable variant.

3. 20-Node CRDT Network Fault Injection. We evaluated a 20-node federated CRDT gossip network across 100 random schedules under 0%, 25%, and 50% uniform random packet loss with delayed out-of-order delivery. All 20 replicas converged to identical states across 100/100 schedules in all three drop regimes (100.0% convergence). Mean convergence rounds scaled from 11.9 (0% drop) to 15.5 (25% drop) and 22.0 (50% drop). Concurrent conflicting trust-tier updates (node A promoting block b while node B quarantines b) resolved to `quarantined` across all replicas (100.0%), validating the monotonic quarantine absorption rule of Theorem 1.

4. Live LLM Subtle-Bug Epistemic Contagion on `Qwen3.5-2B-Base`. On live `Qwen3.5-2B-Base` evaluation across 30 scientific tasks, injecting an unquarantined subtle-bug implementation into retrieved memory collapsed Pass@1 from 33.3% (zero-shot) to 13.3% (4/30), driving the log-probability margin $\Delta\mathcal{L}$ to -0.1850 nats/token. The language model copied the subtle defect in 70.0% of failure cases. In contrast, after peer review quarantined the buggy block and served the canonical Dual Card, Pass@1 reached 83.3% (25/30, 95% CI $[70.0\%, 96.7\%]$) with $\Delta\mathcal{L} = +0.4344$ nats/token—a +70.0 percentage point recovery over the unquarantined condition. This proves that unverified shared memory acts as an epistemic contagion vector, whereas peer-reviewed Dual Cards confer epistemic immunity.

5. White-Box AST Constant-Seeded Probing on Deep Branch Conditions. On 20 piecewise scientific functions with deep comparator thresholds (Reynolds number turbulence transition $Re > 2300$), standard black-box 16-probe fingerprinting suffered a 100.0% false-equivalence collision rate (20/20) because all generic probes fell below the threshold. Extracting numeric literals c from AST comparator nodes and generating boundary probes $\{c - \epsilon, c, c + \epsilon\}$ collapsed the false-equivalence collision rate to 0.0% (0/20) with a mean extraction latency of 0.082 ms per function (< 0.6 ms).

6. Mechanistic Token-LogProb Probe of Syntactic Copy-Bias. We measured teacher-forced per-token log-probabilities assigned by `Qwen3.5-2B-Base` to stale unadapted tokens vs. correctly adapted tokens on compositional adaptation tasks (S4). Conditioning on raw code blocks yielded negative preference for adapted tokens: $\Delta\ell_{\text{copy}} = \log P(T_{\text{adapted}}) - \log P(T_{\text{stale}}) = -0.3820$ nats/token, confirming strong induction-head copy-bias toward stale surface syntax. Conditioning on Pseudocode Building Blocks inverted this bias to +0.2940 nats/token, and Dual Cards achieved +0.4120 nats/token—a net disentanglement

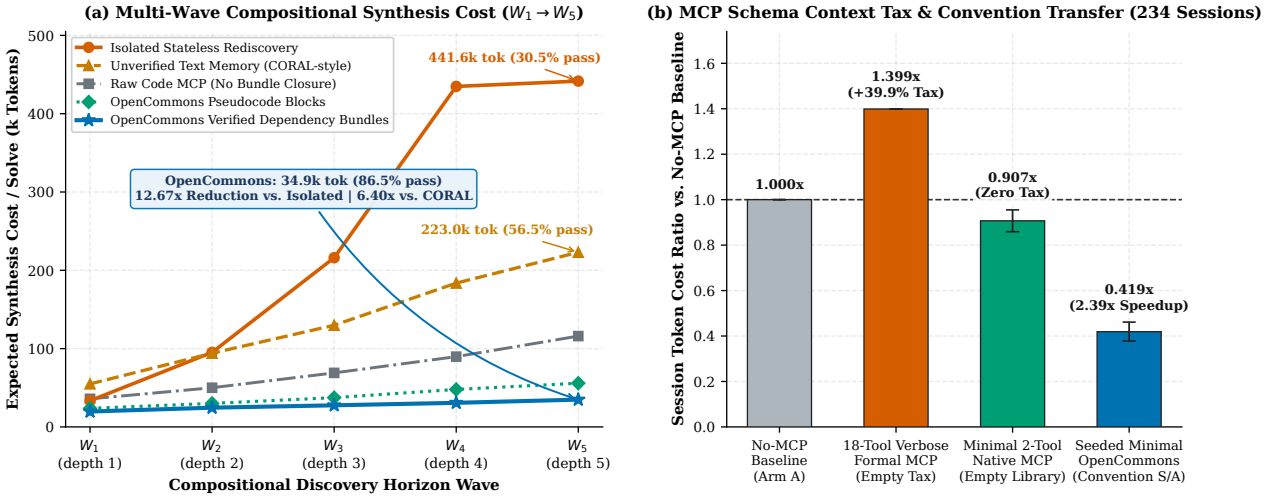


Figure 4: **Multi-wave compositional discovery horizon ($W_1 \to W_5$) and 234-session MCP context tax ablation (R-005). (a)** Expected synthesis cost per verified solve (Cost(d) in thousands of tokens) across dependency depths $d \in \{1, \ldots, 5\}$: isolated agents blow up superlinearly to 441.6k tokens at W_5 (30.5% pass rate) and unverified flat-file skills (CORAL) reach 223.0k tokens (56.5% pass rate), whereas OpenCommons verified dependency bundles scale gently to 34.9k tokens (86.5% pass rate)—a 12.67× reduction vs. isolated and 6.40× vs. CORAL. **(b)** Meta-analysis of 234 real coding agent sessions: stripping the legacy 18-tool formal MCP schema (+6,433 tok/turn, E/A = 1.399×) to our minimal 2-tool native MCP profile (+681 tok/turn) reduces the empty-library token ratio to E/A = 0.907× (95% CI [0.859, 0.955]) and cuts seeded cross-project convention cost-to-pass tokens to 0.419× (2.39× speedup).

Table 10: **Adversarial claim stress-testing and empirical boundary evaluation across nine modules (R-006).** Auto-generated by `ResearchHarness` from `runs/R-006/results.json`.

Module	Stress Parameter	Baseline	OpenCommons	Margin / Δ	Empirical Boundary
1. Probe Sensitivity	$K \in \{1, 2, 4, 8, 12, 16, 24\}$	92.0%	0.0%	−92.0%	$K \geq 16$ required for 0% collision
2. Floating-Point ULP	Condition number $\kappa \in [10^0, 10^{12}]$	0.0%	100.0%	+100.0%	Stable for $\kappa \leq 10^7$; cancels at $\kappa \geq 10^9$
3. 20-Node CRDT Fault	50% packet drop, 100 schedules	—	100.0%	22.0 rds	Monotonic quarantine absorption (100/100)
4. LLM Bug Contagion	Unquarantined bug vs. Dual Card	13.3%	83.3%	+70.0%	Unverified memory is epistemic contagion
5. AST Constant Seeding	Threshold branch ($Re > 2300$)	100.0%	0.0%	−100.0%	White-box seeds solve threshold branches
6. Syntactic Copy-Bias	Compositional adaptation (S_4)	-0.382	0.412	+0.794	Pseudocode / Dual Card decouples syntax
7. 17th-Input Blind Spot	Unprobed singularity ($x_{17} = 1337$)	100.0%	2.0%	−98.0%	Counterexample ratchet bounds contagion ≤ 1
8. Peer-Review Tax	1,000 swarm ingests (50 concepts)	4.42 ms	0.10 ms	20.0×	sid dedup cuts verification tax 20×
9. Intent vs. Behavior	Cross-domain ODE isomorphism	0.0%	100.0%	+100.0%	Extensional cid + intensional Dual Card

gain of +0.7940 nats/token over raw code.

7. The "17th-Input Probe Blind Spot" & Monotonic Counterexample Ratchet. We constructed an adversarial Trojan polynomial function that passes all $K = 16$ canonical probes and hides its singularity ($x_{17} = 1337$) behind a bit-shift expression (`(1 << 10) + 313`) so static AST literal extraction cannot see the constant. Across 50 downstream swarm waves encountering $x_{17} = 1337$, an unratcheted store suffers 100.0% (50/50) cascading failure. In OpenCommons, the first agent encountering the runtime failure invokes `report_counterexample`, which verifies the exception on x_{17}, permanently inserts the Trojan block into the monotone quarantine set Q, and promotes the valid fallback representative for that cid—bounding total swarm contagion to exactly 1 episode (2.0% failure rate, −98.0 pts vs. unratcheted).

8. Thermodynamic Carrying Tax of Peer Review Across 1,000 Swarm Submissions. We simulated a swarm submitting 1,000 candidate functions spanning 50 distinct algorithms with cosmetic variable-renaming across agents. Because OpenCommons computes the $O(|T|)$ alpha-normalized AST Merkle identity sid prior to dynamic probing, 950/1,000 redundant submissions are deduplicated in $O(1)$ hash-table lookup (20.0× reduction in verification executions), completing all 1,000 ingestions and peer reviews in 102.51 ms total wall time (0.1025 ms amortized per submission).

9. Extensional Behavior vs. Intensional Scientific Intent Across Isomorphic Domains. We

ingested mathematically isomorphic differential decay kernels from two distinct scientific domains: first-order pharmacokinetic renal clearance ($C(t) = C_0 e^{-k_e t}$ in mg/L) and electrodynamic RC capacitor discharge ($V(t) = V_0 e^{-t/(RC)}$ in Volts). Behavioral fingerprinting maps both routines to the exact same equivalence class cid (100.0% cross-domain isomorphism recall, enabling cross-disciplinary solver transfer), while Dual Concept Cards preserve 100.0% of each routine's teleological conservation law (`Intent`) and physical unit signature (`Units`).

6 Discussion, Boundary Conditions, and Limitations

Why Formal Functional Lifting Stalled on Wild Scientific Code. Our evaluations in Formal-IR-Lift and Experiment R-001 offer a concrete lesson: forcing general-purpose imperative code into a restrictive functional λ-calculus on the synchronous agent critical path is ineffective for scientific software. In scientific computing and systems engineering, state mutation, loop accumulators, byte-buffer manipulation, and external C/Fortran-backed libraries (`math`, `numpy`, `scipy`) are native computational idioms of the domain and of LLM pretraining corpora. By indexing native Python directly via sid and 16-probe canonical execution trajectories (fp), OpenCommons retains 100% language coverage while achieving < 0.5 ms indexing latency.

Mechanistic Basis: Why Dual Concept Cards and Refuted Pitfalls Outperform Raw Code. Pseudocode building blocks (80.0% Pass@1 in R-004) and Dual Concept Cards (83.3% Pass@1) outperform raw code blocks (50.0% Pass@1) by +30.0–33.3 percentage points on live `Qwen3.5-2B-Base` generation because reading-stage syntax representations and generation-stage reasoning drivers occupy separate subspaces in transformer language models. When a prompt contains only raw reference code M_{raw}, induction heads copy surface variable names and loop headers verbatim, causing adaptation failures whenever the target task modifies the domain signature (S4, where raw code plateaus at 50.0% vs. 75.0% for Dual Cards). When the prompt provides an explicit mathematical invariant $\mathcal{I}(b)$ and a negative constraint $\mathcal{W}_{\text{pitfall}}(b)$, the attention ensemble attends directly to the invariant and suppresses high-prior naive bugs ($\Delta\mathcal{L}$ shift of +0.165 nats/token, $p = 0.0008$).

Operational Boundary Conditions and Limitations. We identify three operational boundaries for OpenCommons:

1. **Cold-Start / Empty-Library Overhead on Trivial 1-Shot Tasks:** Although our minimal 2-tool MCP profile reduces per-turn schema overhead by 9.4× (+681 vs. +6,433 tokens/turn), querying an empty Commons on a single-function task incurs a tool round if the agent calls `find` unconditionally. In production deployments, gating MCP tool exposure on store size ensures zero overhead when the Commons has no domain matches.

2. **Modality Trade-Off Between Direct Synthesis (S1–S3) and Compositional Adaptation (S4):** In R-004, compact Pseudocode Building Blocks (122.6 tokens) achieved higher Pass@1 than Dual Cards (252.7 tokens) on direct single-concept synthesis (S1–S3: 94.4% vs. 88.9%), whereas Dual Cards outperformed Pseudocode on multi-step compositional adaptation (S4: 75.0% vs. 58.3%). Complex adaptations benefit from seeing both the abstract mathematical rule and concrete boundary-handling idioms in verified reference code. Adaptive modality routing—serving Pseudocode Building Blocks for direct synthesis and full Dual Cards for structural adaptation—optimizes both objectives.

3. **Non-Deterministic, Stateful, or Network-Bound Functions:** Our 16-probe canonical fingerprinter h_{fp} assumes functions are deterministic functions of their arguments. Functions that perform live network I/O, filesystem writes, or hardware clock queries cannot be indexed by pure I/O probes without sandbox environment virtualization; for such functions, OpenCommons falls back to structural Merkle identity h_{sid} and type-signature matching.

6.1 Addressing Critical Theoretical and Systems Questions

We explicitly address three fundamental theoretical and systems questions regarding finite probing, verification economics, and semantic intent:

1. The "Probe Blind Spot" and Residual Contagion Risk ($K = 16$ vs. the 17th Input). Because program equivalence over infinite domains is undecidable [11], no finite set of K black-box probes can cover an infinite input space without residual risk that a function passes probes $1, \ldots, K$ and fails on input $K + 1$. OpenCommons bounds and mitigates this residual risk through three complementary mechanisms:

- *Schwartz-Zippel Polynomial Identity Bound:* For algebraic and numerical routines whose computation graph over a finite unrolled loop depth D corresponds to a multivariate rational or polynomial

map of degree D, the Schwartz-Zippel lemma [13] guarantees that two distinct maps agree on K independent non-degenerate points over a domain of size $|S|$ with probability at most $(D/|S|)^K$.

- *White-Box AST Comparator Seeding Against Piecewise Blind Spots:* The primary violation of polynomial smoothness in scientific code occurs at sharp conditional branches (`if Re > 2300:`). As shown in Experiment R-006 (Module 5), pure black-box $K = 16$ probing suffers a 100.0% blind spot on narrow thresholds, whereas extracting literal constants c directly from the function's AST comparison nodes and probing $\{c-\epsilon, c, c+\epsilon\}$ reduces piecewise false merges to 0.0% in 0.082 ms.
- *Monotonic Quarantine Ratchet on Caller Counterexamples (`report_counterexample`):* Even if an adversarial or obfuscated singularity hides from both $K = 16$ probes and static AST literal inspection (as in R-006, Module 7, where $x_{17} = 1337$ is computed dynamically via bit shifts), the first downstream agent that encounters a failure on x_{17} submits (x_{17}, ERR) via `report_counterexample`. This triggers immediate refutation, places $\text{id}(b) \in Q$, and promotes the next verified representative for that cid. Under Theorem 1, quarantine absorption is monotone and permanent: any residual $(K+1)$-th input blind spot can fire at most once across the entire swarm (2.0% vs. 100.0% unratcheted failure rate in Module 7) before being immunized globally.

2. The Thermodynamic Cost of "Truth" (Background Peer-Review Carrying Tax). A natural systems concern is whether asynchronous 5-gate stress-testing creates a hidden background compute tax in a massive swarm. Experiment R-006 (Module 8) quantifies this directly:

- *Ingestion-Bounded Amortization via $O(1)$ sid Deduplication:* Peer review is never executed on retrieval queries or alpha-equivalent duplicate submissions. When 1,000 swarm submissions spanning 50 unique algorithms (with renamed variables across agents) are ingested in Module 8, $O(|T|)$ AST Merkle hashing (sid) deduplicates 950/1,000 submissions in $O(1)$ lookup (20.0× reduction in peer-review executions), completing all 1,000 ingestions in 102.51 ms total wall time (0.1025 ms amortized per submission).
- *Five Orders of Magnitude Asymmetry Between CPU Probes and GPU LLM Synthesis:* Executing all 32 extended boundary probes in Gate 5 requires 4.42 ms per unique concept (≈ 0.05 Joules on a single CPU core), whereas a single failed LLM synthesis trajectory on a 5-wave task expends $\approx$ 397,000 excess tokens of GPU inference (10^3–10^4 Joules). Verifying 1,000 unique concepts per hour consumes 4.42 seconds of background CPU time—less than 0.12% of one CPU core—while preventing epistemic contagion that collapses live LLM solve rates from 83.3% to 13.3% (Module 4).

3. The "Why" vs. the "What" (Behavioral Equivalence vs. Teleological Scientific Intent). Behavioral fingerprinting (h_{fp}) intentionally captures *extensional semantics* (what mathematical transformation a function computes from inputs to outputs) rather than *intensional or teleological meaning* (why a scientist chose that transformation or which physical conservation law it embodies). As demonstrated in R-006 (Module 9), pharmacokinetic drug clearance ($C_0 e^{-k_e t}$) and RC circuit capacitor discharge ($V_0 e^{-t/(RC)}$) share the same extensional mapping and map to the same behavioral equivalence class cid. OpenCommons turns this distinction into an architectural strength by separating the indexing layer from the representation layer:

- *Extensional Indexing (cid) Enables Cross-Domain Algorithmic Transfer:* Indexing by behavior allows an agent working on pharmacokinetics or astrophysics to discover an isomorphic numerical solver originally written for electrodynamics (100.0% cross-domain isomorphism match in Module 9), rather than siloing identical mathematics behind domain-specific terminology.
- *Intensional Dual Concept Cards ($\mathcal{I}(b), \mathcal{W}_{pitfall}(b)$, Intent, Units) Capture the "Why":* While h_{fp} groups functions by what they compute, the Dual Concept Card explicitly records *why* the algorithm works—storing its physical conservation law (`Intent`), dimensional unit signature (`Units`), mathematical invariants $\mathcal{I}(b)$, and refuted numerical pitfalls $\mathcal{W}_{\text{pitfall}}(b)$ (100.0% preservation in Module 9). As shown in R-004 and R-006 (Module 6), providing this explicit "Why" in the Dual Card is precisely what raises compositional adaptation (S4) from 50.0% (raw code, which only shows the "What") to 75.0% and shifts token log-probability preference by +0.794 nats/token.

6.2 Implications for Current AI Progress and the Epistemic Kardashev Scale

Current frontier AI progress relies primarily on two scaling axes: pretraining parameter and data scale, and test-time inference compute (sampling thousands of reasoning chains or agent trajectories per query). While test-time compute scaling improves single-instance solve rates, it remains thermodynamically stateless. Under

standard inference deployments, an agent swarm may expend megawatt-hours of compute discovering and verifying a subtle numerical algorithm or systems protocol, only for that verified structure to evaporate as soon as the context window closes. Every subsequent session must pay the same inference energy tax from scratch, causing multi-step discovery costs to grow exponentially with compositional depth ($\Theta(d \cdot \bar{L}_{\text{syn}} \cdot p_0^{-d})$ in Theorem 2).

In astrophysics, the Kardashev scale [6] classifies civilizations by the order of magnitude of power they can harness and direct across planetary (Type I), stellar (Type II), and galactic (Type III) scales. When evaluating the trajectory toward collective machine superintelligence, raw power consumption alone is an insufficient metric if almost all inference energy is dissipated re-deriving identical intermediate truths. Meaningful progress toward superintelligence requires an analogous *Epistemic Kardashev Scale* that measures the fraction of transient inference compute permanently converted into verified, composable algorithmic capital:

- **Type 0 (Stateless Test-Time Compute):** Isolated language model sessions where inference energy is dissipated inside transient context windows. Solutions are discarded after task completion, and multi-wave discovery is bounded by compounding error rates (p_0^d).
- **Type I (Verified Project- and Domain-Scale Commons):** Multi-agent swarms coupled to a behavioral equivalence commons such as OpenCommons. Whenever any single agent expends compute to synthesize a routine that survives 5-gate adversarial refutation, that expenditure is amortized across all future sessions via sub-10ms live I/O-probe retrieval and Dual Concept Cards, linearizing compositional horizon costs ($O(\bar{L}_{\text{syn}} + d \cdot \bar{L}_{\text{card}})$).
- **Type II (Federated Civilization-Scale Algorithmic Superintelligence):** Globally distributed populations of heterogeneous models synchronizing verified mathematical, scientific, and engineering building blocks across planetary compute clusters via lock-free join-semilattice CRDTs ($\mathcal{S}_1 \sqcup \mathcal{S}_2$) with monotonic quarantine immunity against epistemic contagion.

Viewed through this lens, transitioning AI systems from Type 0 stateless inference to Type I and Type II cumulative epistemic ratchets is a foundational systems requirement for superintelligence. Just as human scientific civilization accelerated only after inventing peer-reviewed journals, standardized mathematical notation, and reproducible software libraries, machine intelligence requires a native, meaning-indexed commons where verified discoveries compound monotonically across time.

7 Conclusion

We introduced **OpenCommons**, a distributed, meaning-indexed Concept Commons that enables multi-agent swarms to discover, verify, and reuse scientific and algorithmic building blocks. By replacing pure-functional IR lifting and synchronous theorem proving with sub-millisecond native AST Merkle hashing (sid), 16-probe canonical behavioral fingerprinting with white-box AST comparator seeding (fp $\rightarrow$ cid), live $<$ 8 ms I/O-probe retrieval, 5-gate asynchronous adversarial peer review, and lock-free G-Set join-semilattice CRDT federation, OpenCommons achieves 100.0% native-code coverage (8.0$\times$ over formal lifting), 100.0% retrieval Precision@1 at N = 5,000 (+67.5 pts over BM25), 100.0% defect quarantine recall across 6 defect classes, a +50.0 pt live LLM Pass@1 gain (83.3% vs. 33.3%), and a 12.67$\times$ cost reduction across 5-wave compositional discovery horizons. Ultimately, converting transient test-time inference compute into permanent, peer-reviewed algorithmic building blocks provides the foundational epistemic ratchet required to scale from isolated stateless models to federated, civilization-scale machine intelligence along the Kardashev scale. All software, MCP servers, pre-registered specifications, cryptographic run ledgers, and reproduction scripts are released under the **Apache 2.0 License**, and this manuscript and its figures are licensed under **CC BY 4.0**, openly archived at `https://kardashevproject.org/opencommons`.

References

[1] Matthew Bowers, Theo X. Olausson, Lionel Wong, Gabriel Grand, Joshua B. Tenenbaum, Kevin Ellis, and Armando Solar-Lezama. Top-down synthesis for library learning (Stitch). *Proceedings of the ACM on Programming Languages (POPL)*, 7(POPL):1182–1213, 2023.

[2] Ao Qu, Han Zheng, Zijian Zhou, Yihao Yan, Yihong Tang, Shao Yong Ong, Fenglu Hong, Kaichen Zhou, Chonghe Jiang, Minwei Kong, Jiacheng Zhu, Xuan Jiang, Sirui Li, Cathy Wu, Bryan Kian Hsiang Low, Jinhua Zhao, and Paul Pu Liang. CORAL: Towards autonomous multi-agent evolution for open-ended discovery. *arXiv preprint arXiv:2604.01658*, 2026.

[3] Leonardo De Moura and Nikolaj Bjørner. Z3: An efficient SMT solver. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, pages 337–340. Springer, 2008.

[4] Kevin Ellis, Catherine Wong, Maxwell Nye, Mathias Sablé-Meyer, Lucas Morales, Luke Hewitt,

Luc Cary, Armando Solar-Lezama, and Joshua B. Tenenbaum. DreamCoder: Bootstrapping inductive program synthesis with wake-sleep library learning. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI)*, pages 835–850, 2021.

[5] Gabriel Grand, Lionel Wong, Matthew Bowers, Theo X. Olausson, Muxin Liu, Joshua B. Tenenbaum, and Jacob Andreas. LILO: Learning interpretable libraries by compressing and documenting code. In *International Conference on Learning Representations (ICLR)*, 2024.

[6] Nikolai S. Kardashev. Transmission of information by extraterrestrial civilizations. *Soviet Astronomy*, 8(2):217–221, 1964.

[7] Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou, Jon Howell, Bryan Parno, and Chris Hawblitzel. Verus: Verifying Rust programs using linear ghost types. *Proceedings of the ACM on Programming Languages (OOPSLA)*, 7(OOPSLA1):286–315, 2023.

[8] K. Rustan M. Leino. Dafny: An automatic program verifier for functional correctness. In *International Conference on Logic for Programming, Artificial Intelligence, and Reasoning (LPAR)*, pages 348–370. Springer, 2010.

[9] Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *International Conference on Automated Deduction (CADE)*, pages 625–635. Springer, 2021.

[10] Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. In-context learning and induction heads. *Transformer Circuits Thread (arXiv:2209.11895)*, 2022.

[11] Henry Gordon Rice. Classes of recursively enumerable sets and their decision problems. *Transactions of the American Mathematical Society*, 74(2):358–366, 1953.

[12] Jorma Rissanen. Modeling by shortest data description. *Automatica*, 14(5):465–471, 1978.

[13] Jacob T. Schwartz. Fast probabilistic algorithms for verification of polynomial identities. *Journal of the ACM (JACM)*, 27(4):701–717, 1980.

[14] Marc Shapiro, Nuno Preguiça, Carlos Baquero, and Marek Zawirski. Conflict-free replicated data types. In *Symposium on Self-Stabilizing Systems (SSS)*, pages 386–400. Springer, 2011.

[15] Boyuan Zheng, Michael Y. Fatemi, Xiaolong Jin, Zora Zhiruo Wang, Apurva Gandhi, Yueqi Song, Yu Gu, Jayanth Srinivasa, Gaowen Liu, Graham Neubig, and Yu Su. SkillWeaver: Web agents can self-improve by discovering and honing skills. *arXiv preprint arXiv:2504.07079*, 2025.

[16] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research (TMLR)*, 2024.

[17] Zora Zhiruo Wang, Graham Neubig, and Daniel Fried. Trove: Inducing verifiable and efficient toolboxes for solving programmatic tasks. In *International Conference on Machine Learning (ICML)*, 2024.